



Time Series Management

Michele Linardi Ph.D.

michele.linardi@orange.fr



Syllabus

- Recap: Time Series Forecasting and Deep Learning Fundamentals
- Introduction to convolutional neural networks (CNN)
- Implementation of the models with the Keras Library

Time series data... quick recap

- A **univariate time series** is a sequence of measurements of the same variable collected over time. Most often, the measurements are made at regular time intervals.



<https://www.kaggle.com/code/anushkaml/walmart-time-series-sales-forecasting>

Time Series Forecasting

- Predict Future (Values) based on the past (Historical Data)

	Date	Observation
History	2018-06-04	60
	2018-06-05	64
	2018-06-06	66
	2018-06-07	65
	2018-06-08	67
	2018-06-09	68
	2018-06-10	70
	2018-06-11	69
	2018-06-12	72
	2018-06-13	?
Future	2018-06-14	?
	2018-06-15	?

Question to ask prior to forecast

- Can we forecast?
 - How well we understand the factor influencing the future?
 - How much data we have?*
- How far in the future (horizon) we want to forecast?
- What technique, model should we use ?

Why Deep Learning models? (1/2)

Deep learning model perform well and better than other methods in many scenarios.

Model	Features	RMSE	MAPE (%)	Run time (s)
FeedForward	Date features + Covariants	10.73	4.20	53.35
AutoARIMA	Baseline	11.00	3.52	20658.95
DeepAR	Date features	11.96	5.99	156.18
FeedForward	Baseline	12.01	3.88	78.47
SeasonalNaive	Baseline	12.47	3.86	1.20
FeedForward	Date features	13.93	3.95	9.37
DeepAR	Baseline	14.01	6.12	856.51
DeepAR	Date features + Covariants	17.79	6.04	781.09
TrivialIdentity	Baseline	18.20	5.09	1.30
NPTS	Baseline	82.45	19.75	25.20

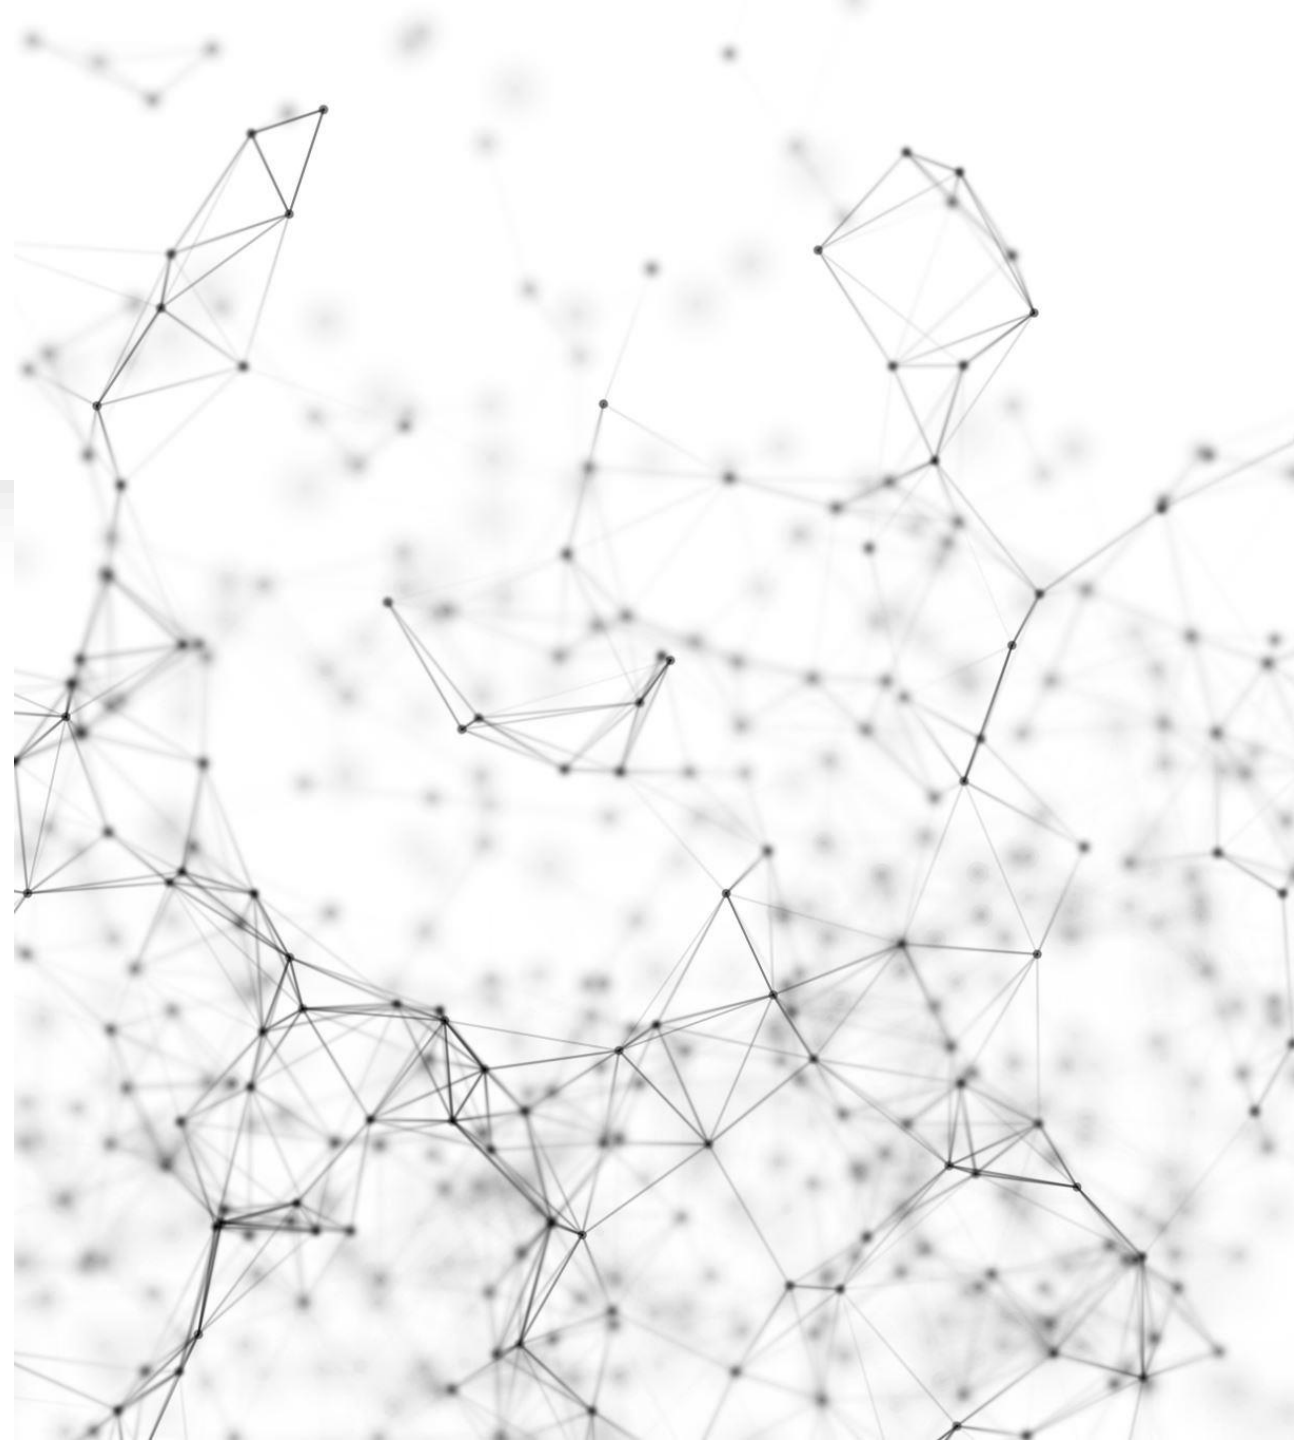
<https://blog.dataiku.com/deep-learning-time-series-forecasting>

Z. Tang, P.A. Fishwik, [Feedforward Neural Nets as Models for Time Series Forecasting](#), November 1993

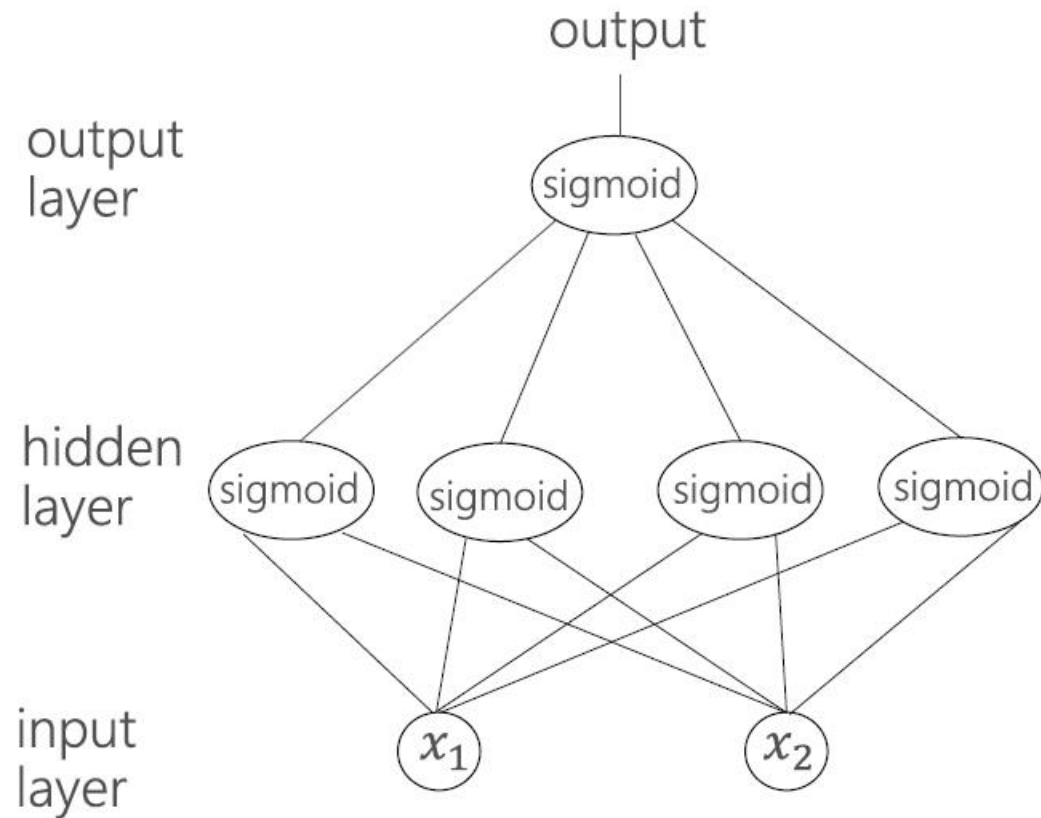
Why Deep Learning models? (2/2)

- Non-parametric
- Flexible and expressive
- Easily inject exogenous features into the model
- Learn from large time series datasets

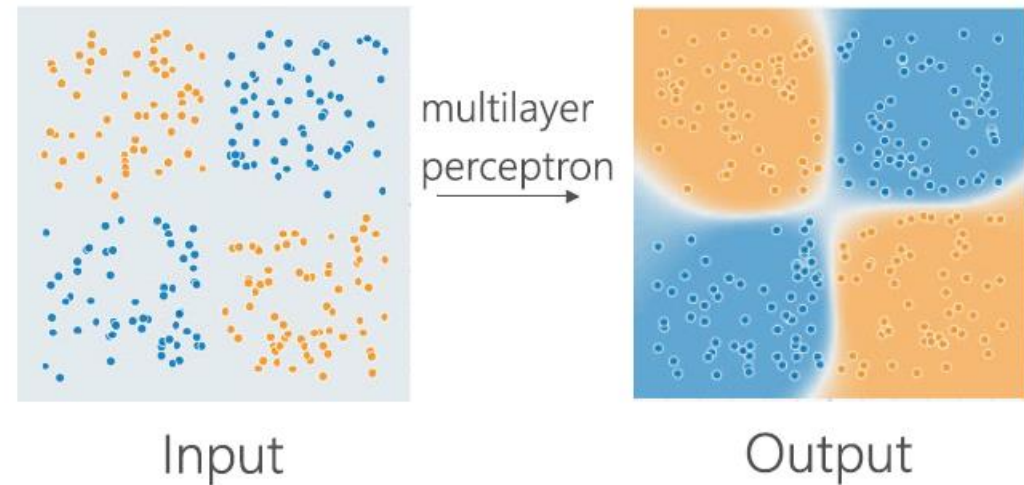
DEEP LEARNING FUNDAMENTALS



Multilayer Perceptron



$$\hat{y} = \begin{cases} 1 & \text{if output} > \text{threshold} \\ 0 & \text{otherwise} \end{cases}$$

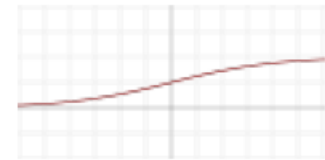


also known as Feed-Forward Neural Network and Deep Neural Network

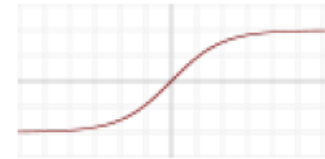
Activation Function

Applied over $x' = w \cdot x + b$

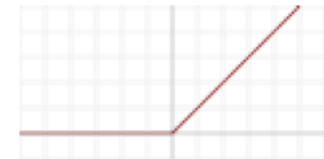
Sigmoid $f(x') = \sigma(x') = \frac{1}{1+e^{-x'}}$



tanh (hyperbolic tangent) $f(x') = \frac{e^{x'} - e^{-x'}}{e^{x'} + e^{-x'}}$




Rectifier linear unit (ReLU) $f(x') = \begin{cases} 0 & \text{for } x' < 0 \\ x' & \text{for } x' \geq 0 \end{cases}$



Training: overview


- Main Idea: Use training example to find proper weights and biases.
- Loss function **$L(\text{Weights}, \text{biases})$** measure the discrepancy between predicted and true values
- We want to optimize weight and biases to reduce the Loss.



Optimization (batch) Gradient descent

Gradient $\nabla L(w) = \left(\frac{\partial L(w)}{\partial w_1}, \frac{\partial L(w)}{\partial w_2}, \dots, \frac{\partial L(w)}{\partial w_d} \right)$ - direction of the maximal increase of $L(w)$

1D example: $\nabla L(w) = \left(\frac{\partial L(w)}{\partial w} \right)$



Optimization algorithm

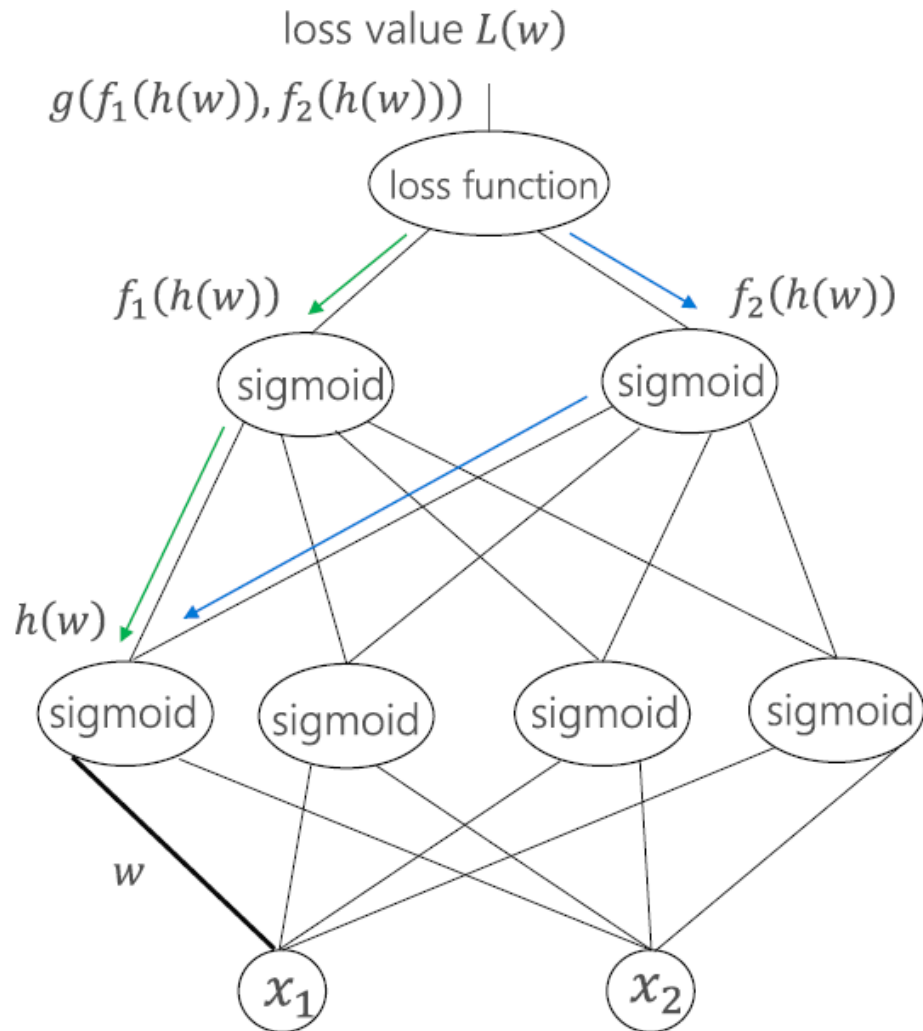
Initialization: $w = w_0$

While stopping criterion not met:

$$w = w - \alpha \cdot \nabla L(w)$$

learning rate

Gradients Computation: Backpropagation



Computation of $L'(w)$

Chain rule

$$L(w) = g(f_1(w))$$

$$L'(w) = g'(f_1(h(w))) \cdot f_1'(h(w)) \cdot h'(w)$$

$$L'(w) = g'(f_2(h(w))) \cdot f_2'(h(w)) \cdot h'(w)$$


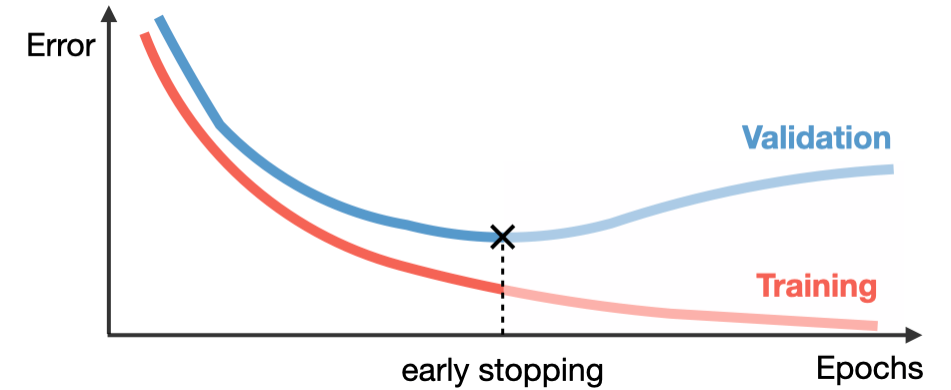
Total derivative

$$L(w) = g(f_1(w), f_2(w))$$

$$L'(w) = L'(w) + L'(w)$$

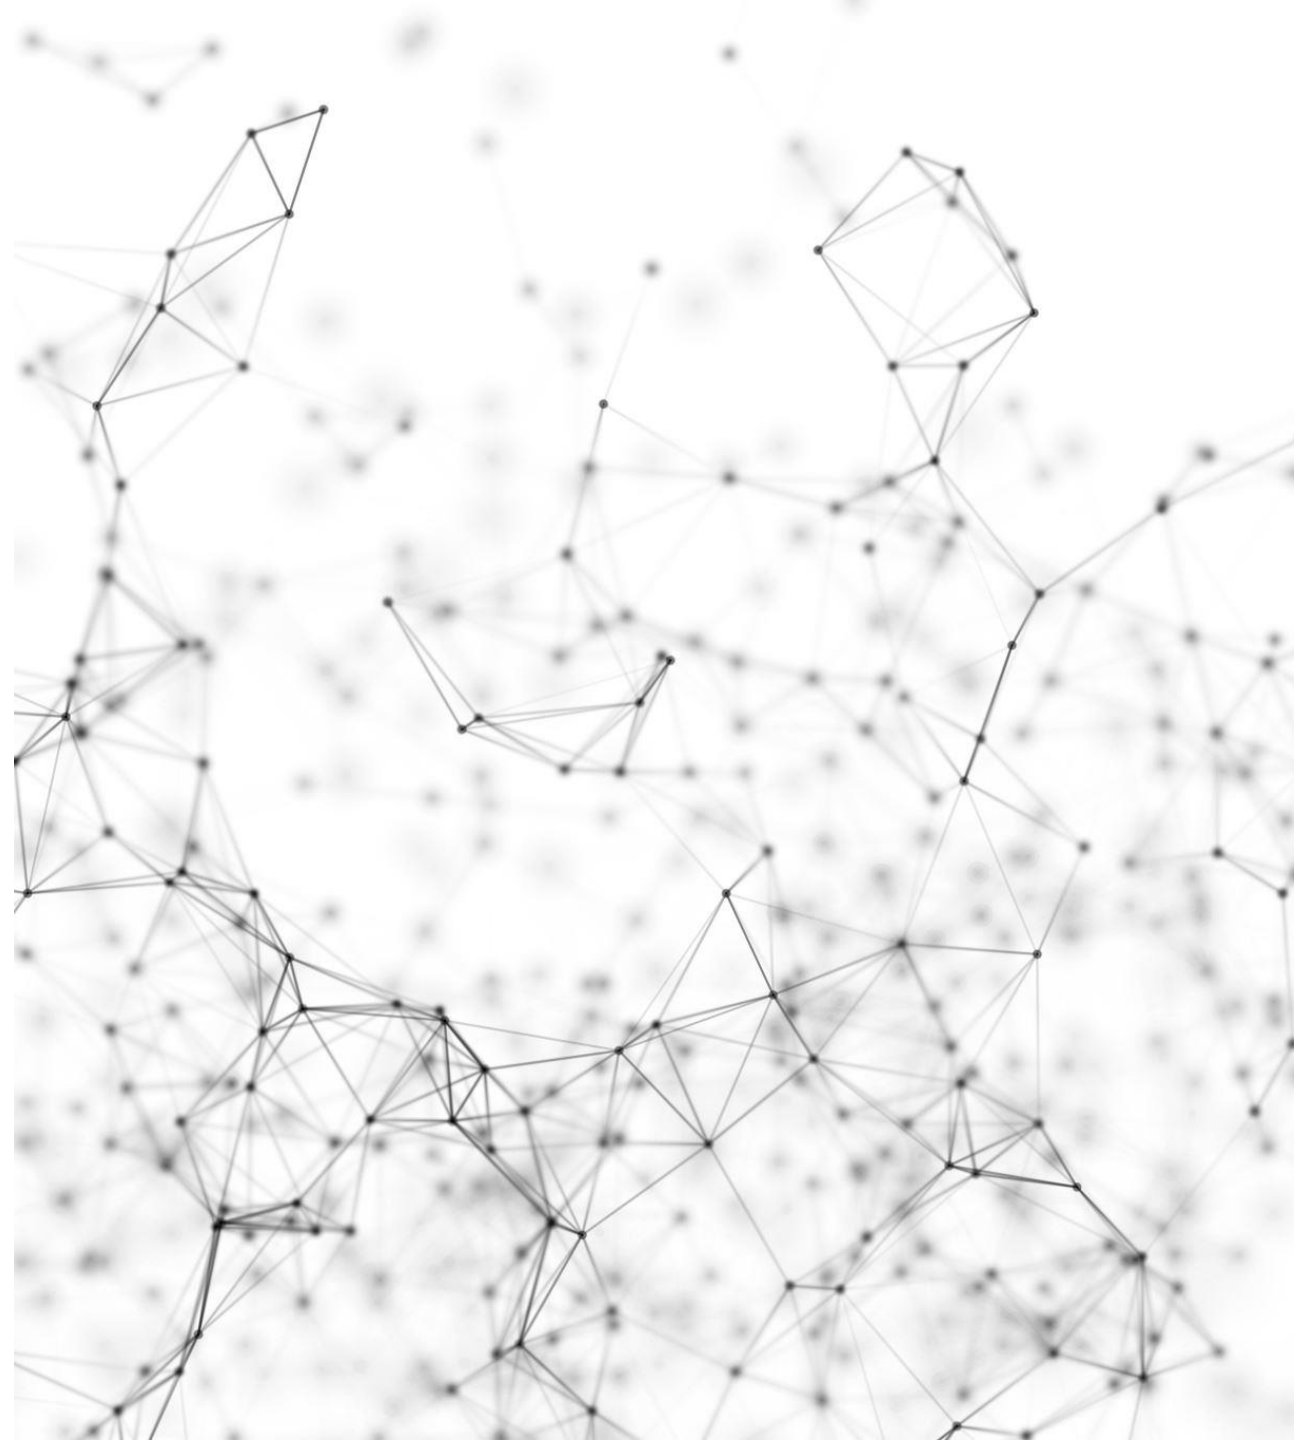
Training improvement strategy

- **Early stopping**
- Tuning hyperparameters
- Initialization
- Weight regularization
- **Dropout**
- **Batch Normalization**
- Learning rate decay



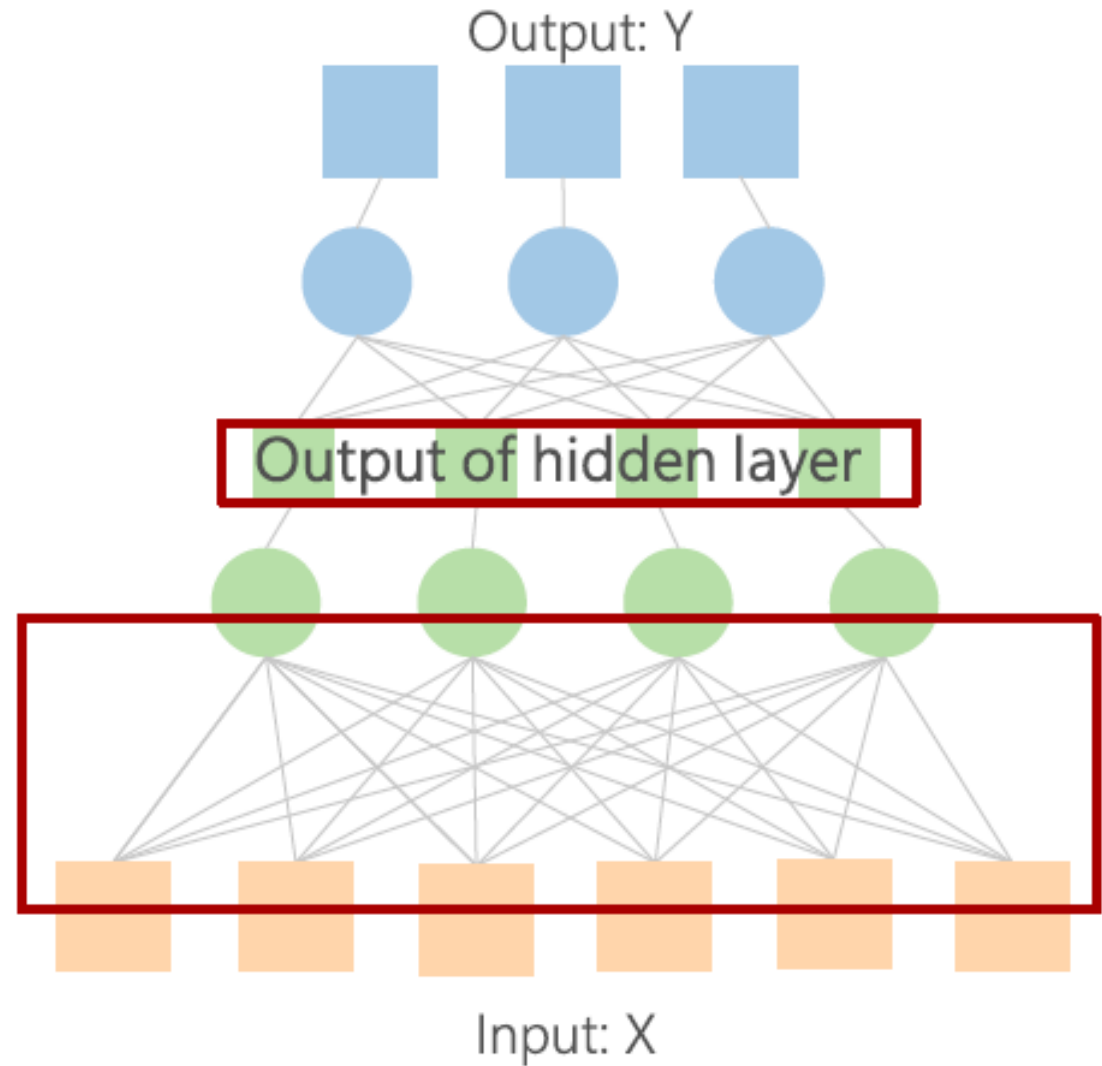
$$x_i \leftarrow \gamma \frac{x_i - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} + \beta$$

Convolutional neural networks (CNN) for Time series Forecasting



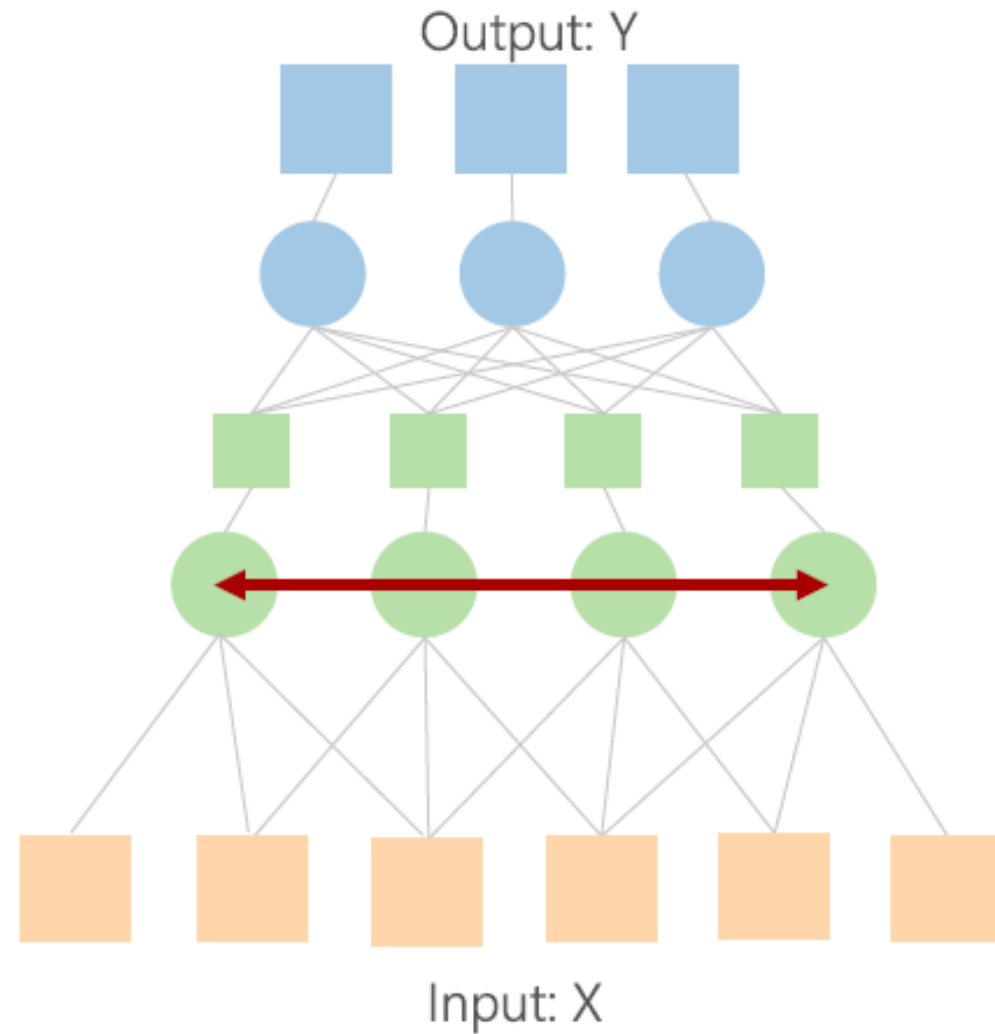
Fully connected VS Convolutional Layer

Fully Connected: Units in hidden layer are connected to every unit in previous layer.



Fully connected VS Convolutional Layer

Convolutional Layer: Units in hidden operate on a field of the input (the weights are shared across the input)



1D Convolution

- Apply a 1D filter to all elements of an input time series.
- Result: SUM of element-wise product

$$\begin{array}{|c|c|c|c|c|c|} \hline 5 & 3 & 2 & 7 & 1 & 6 \\ \hline \end{array} \times \begin{array}{|c|c|c|} \hline 1 & 0 & -1 \\ \hline \end{array} = \begin{array}{|c|c|c|c|} \hline 3 & -4 & 1 & 1 \\ \hline \end{array}$$

$(5 \times 1) + (3 \times 0) + (2 \times -1) = 3$

Input: 1 x 6 Filter: 1 x 3 Output: 1 x 4

1D Convolution

- Filters are trained to detect features in the input sequence.
- Features can be detected regardless of where they appear in the input sequence.

5	3	2	7	1	6
---	---	---	---	---	---

X

w_1	w_2	w_3
-------	-------	-------

=

3	4	1	1
---	---	---	---

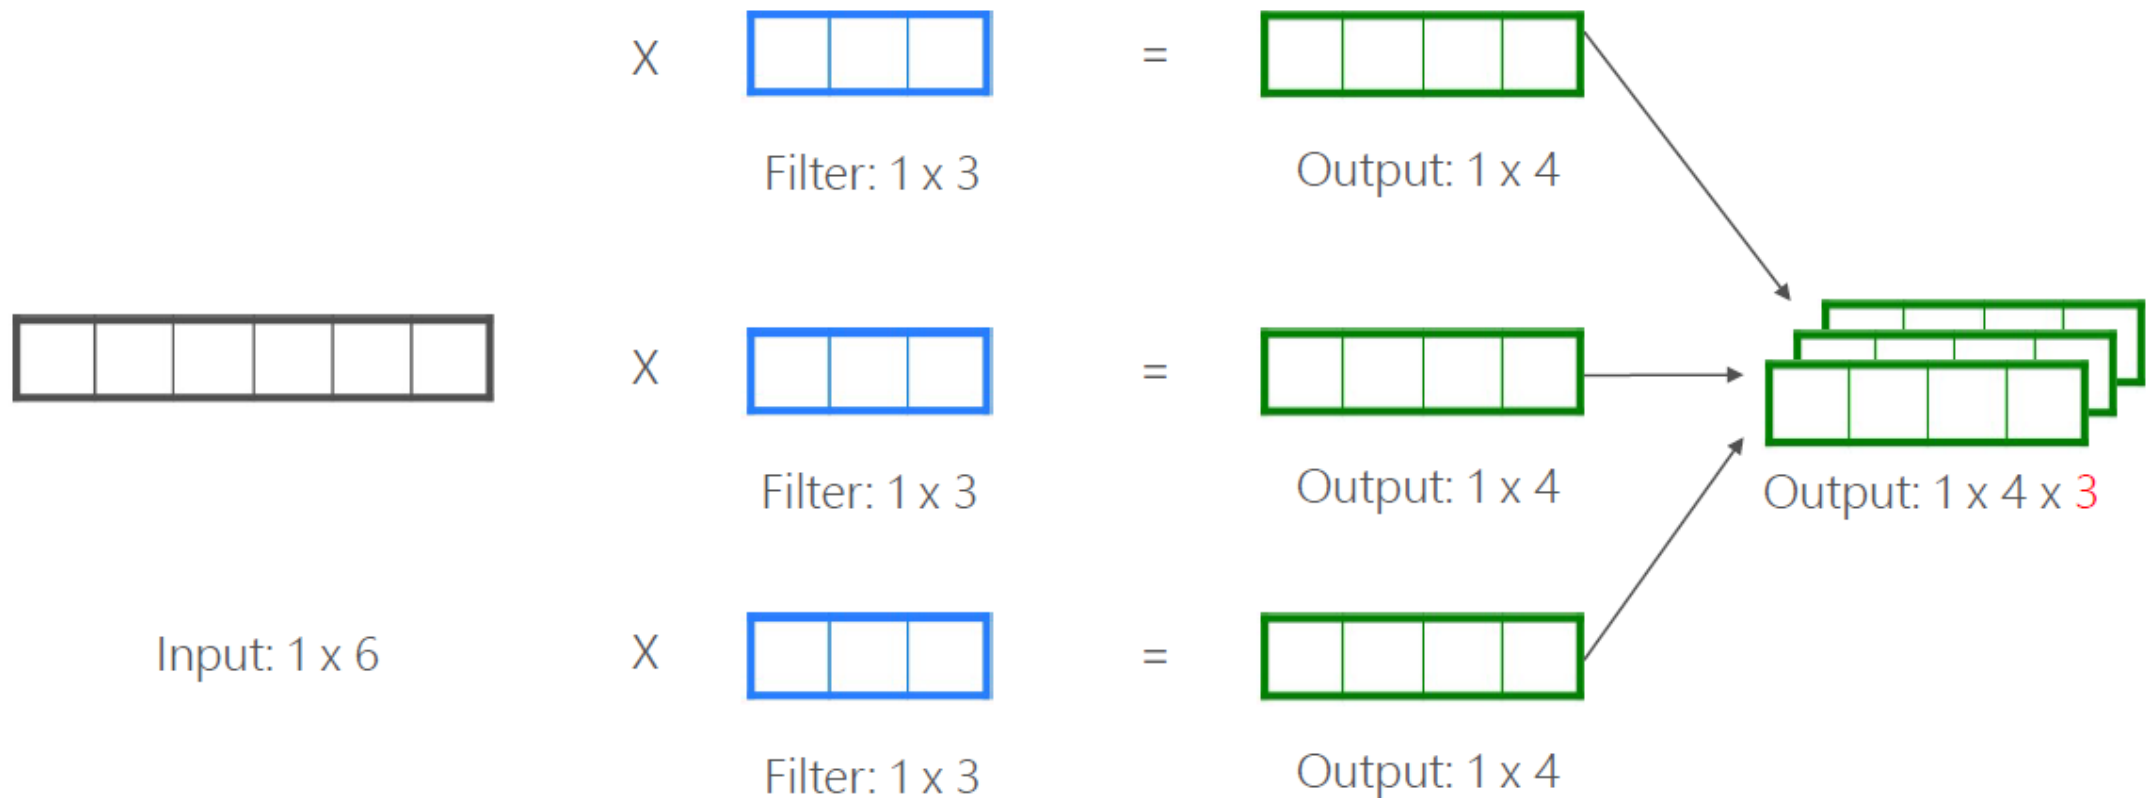
Input: 1 x 6

Filter: 1 x 3

Output: 1 x 4

1D Convolution


- Apply multiple filters to the input data to detect multiple features.



Apply CNN to forecasting time series

At time t :

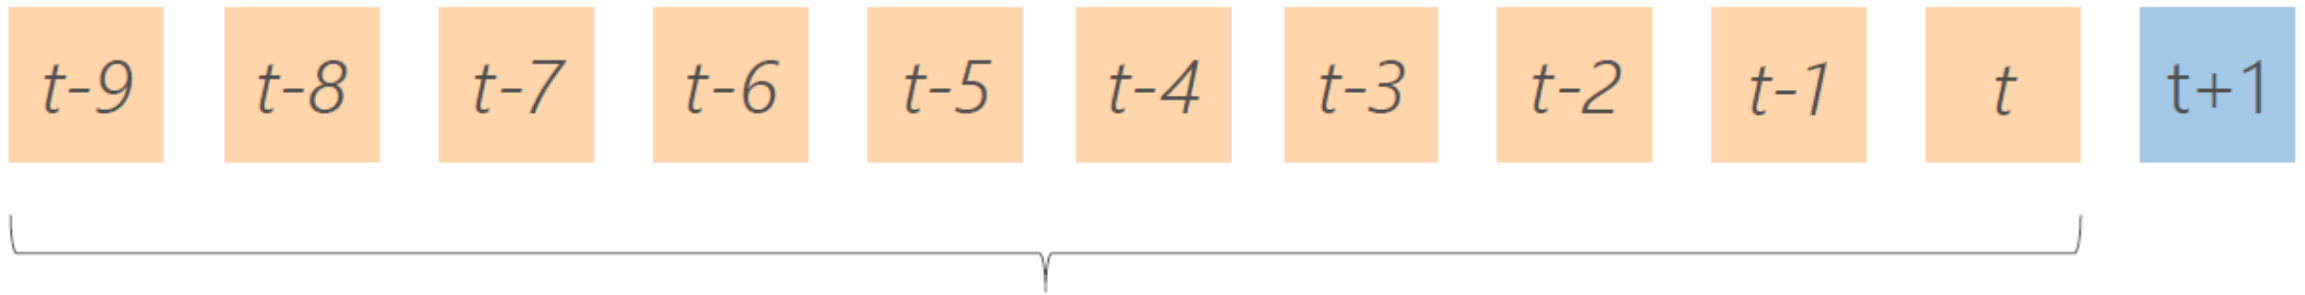
1. predict the value at time $t+1$
2. Conditioned on the previous T values of the time series



Apply CNN to forecasting time series

At time t :

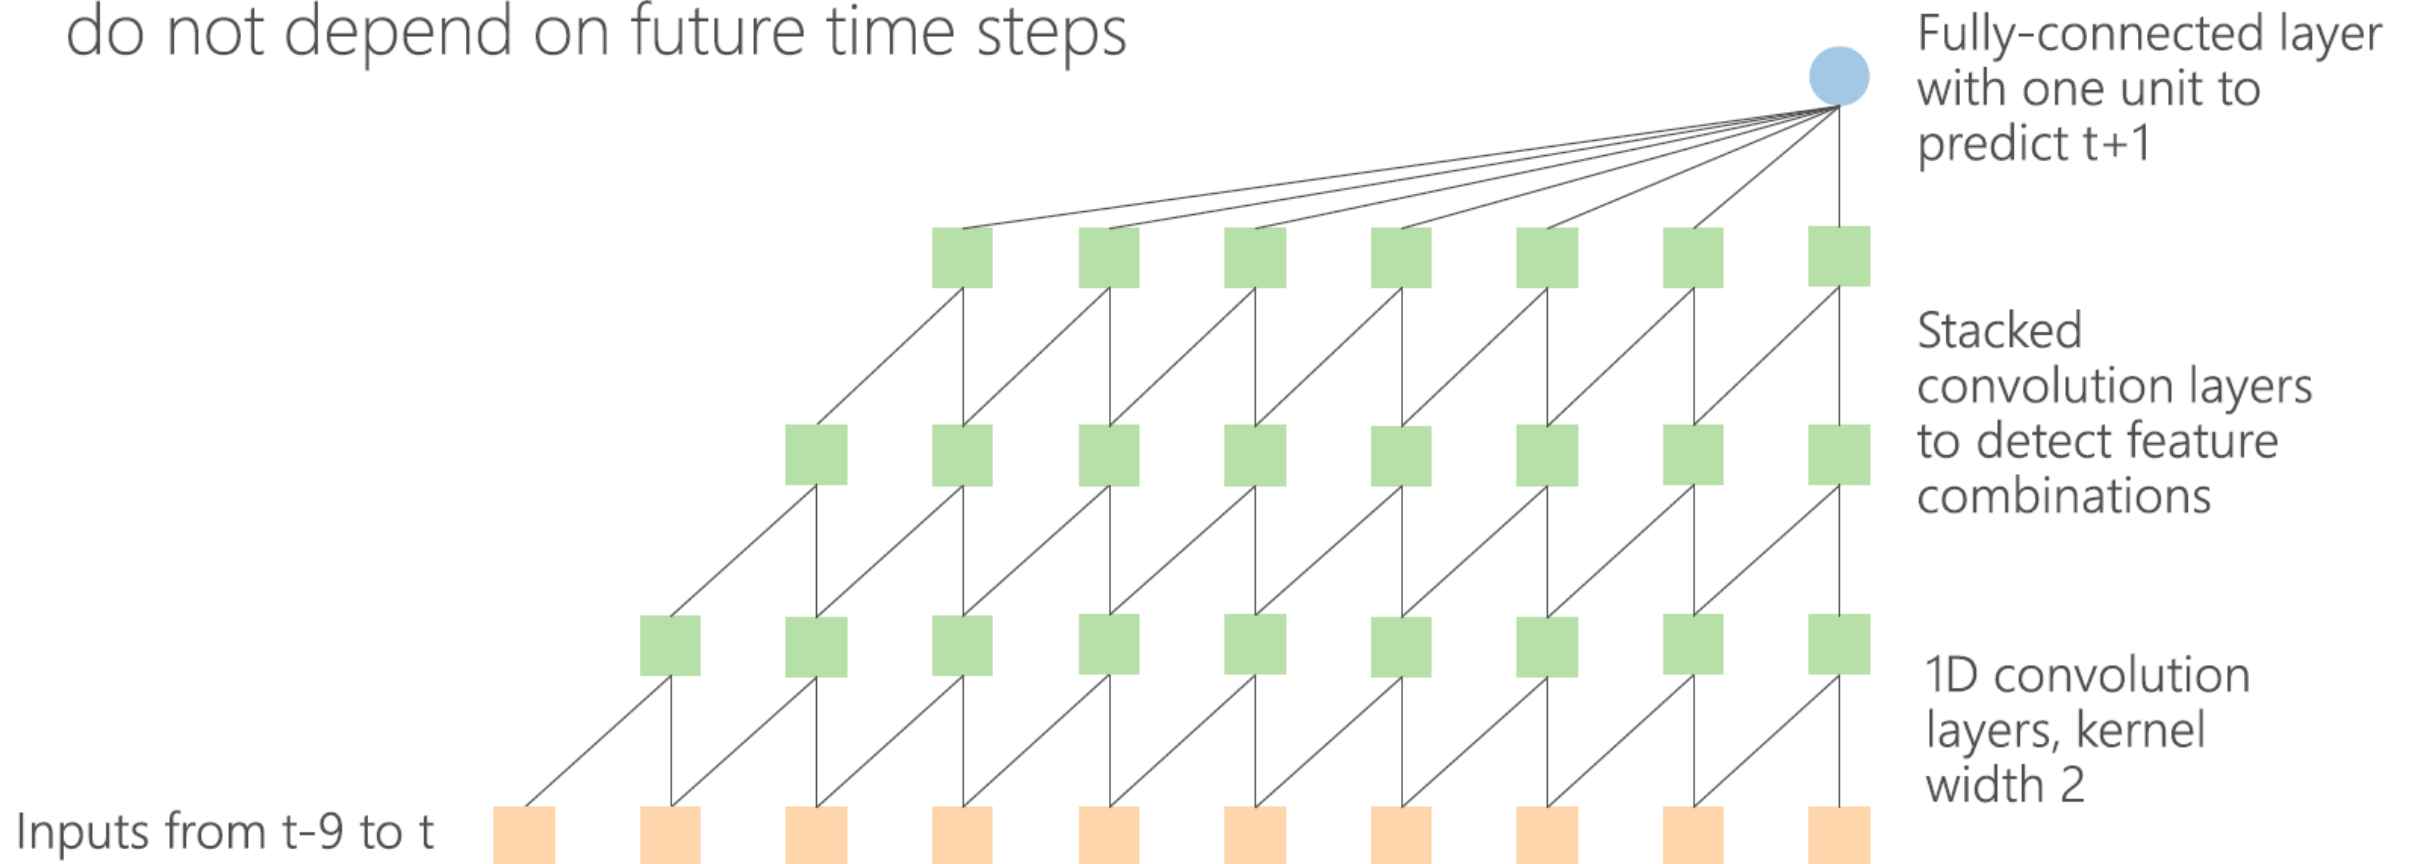
1. predict the value at time $t+1$
2. Conditioned on the previous T values of the time series



$T = 10$ --- Empirically selected


Apply CNN to forecasting time series

- Causal convolutions: the output at each time step do not depend on future time steps



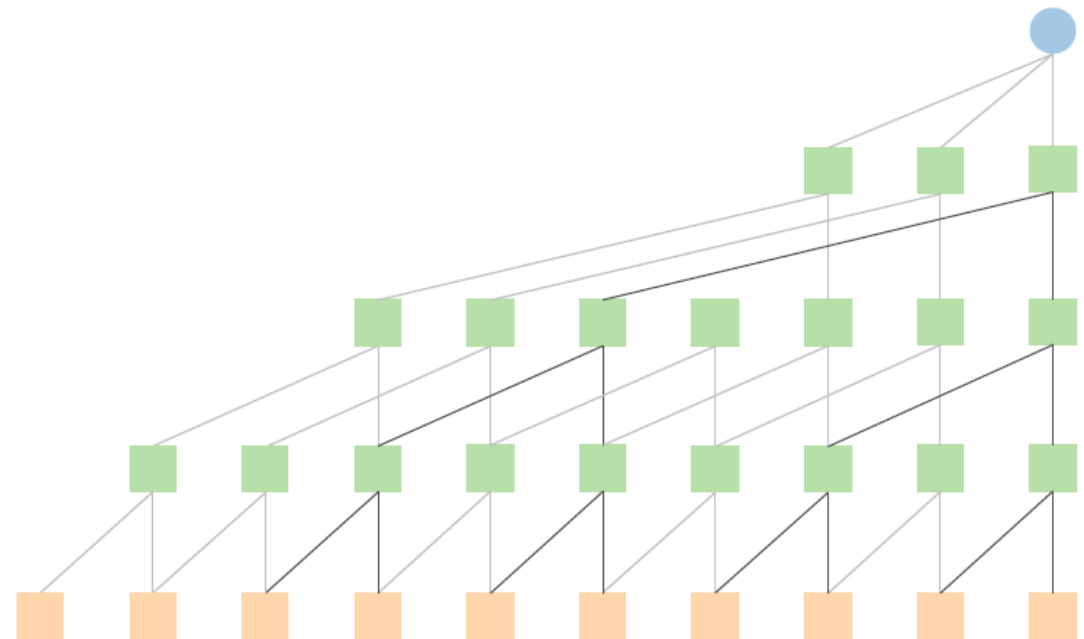
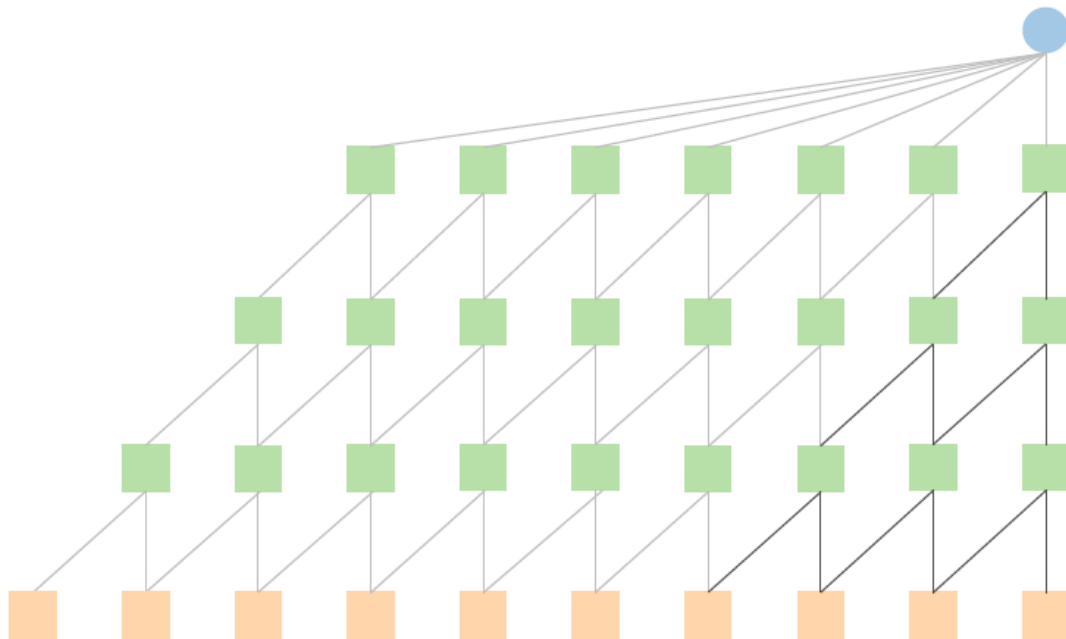
Dilated Convolutions

- Skip output from previous layers



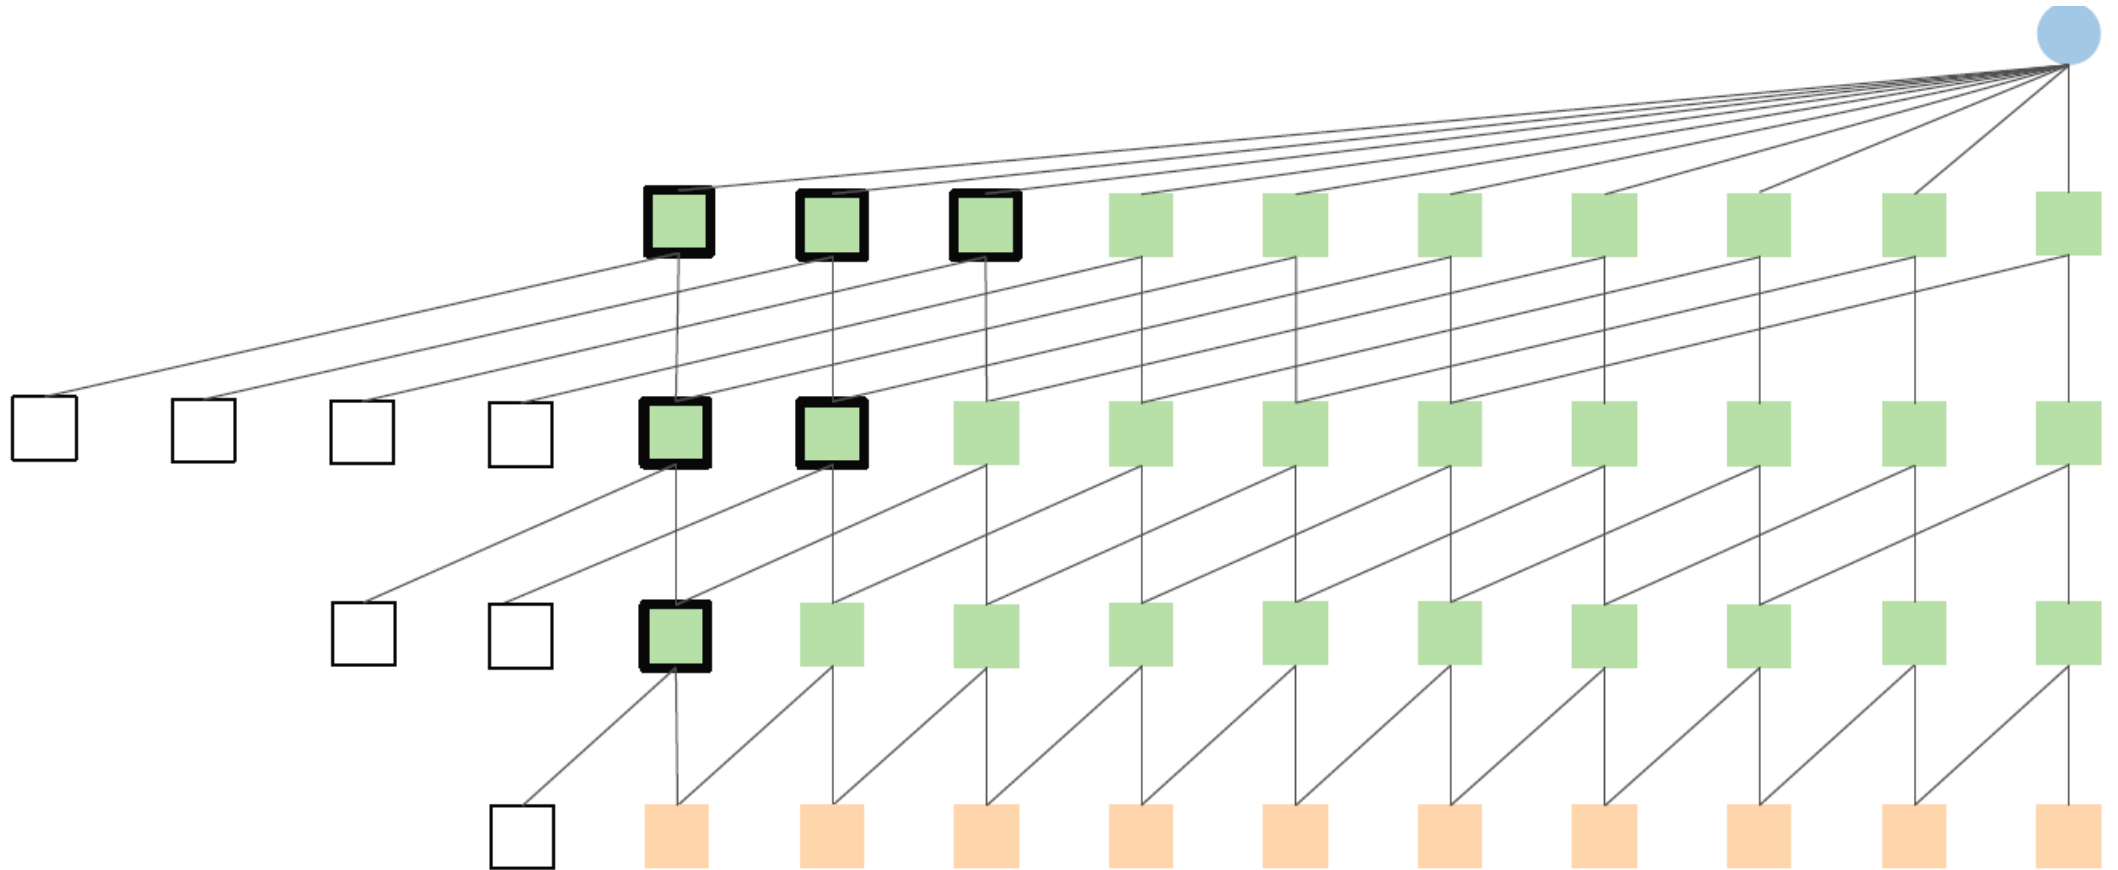
Dilated Convolutions

- Normal convolution => more connections => more weights to train
- Reach information from more distant values in the time series (non contiguous filter)



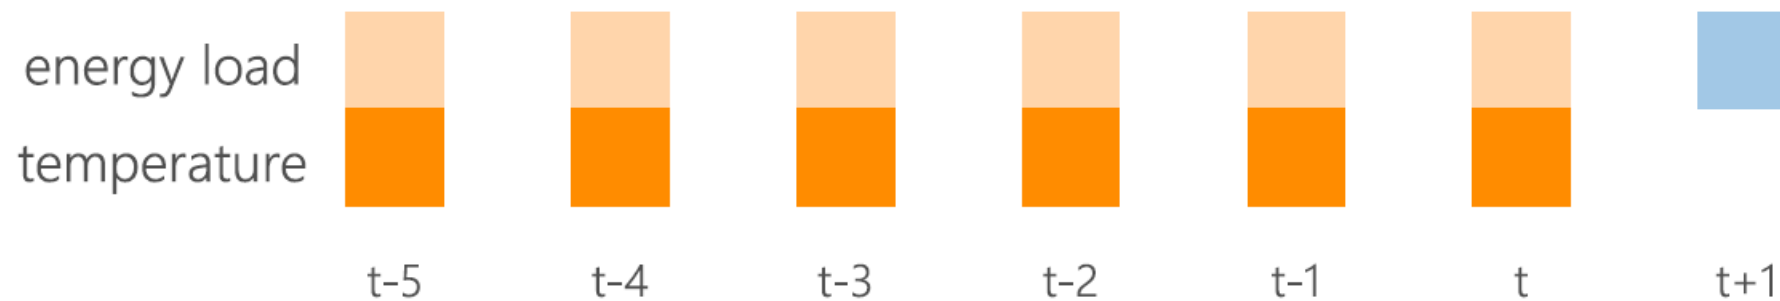
Causal padding

- Padding is necessary to keep the right output dimensions




Multivariate Time Series Case

- Time series with more than 1 variable evolving along time (aka covariates)
- Example

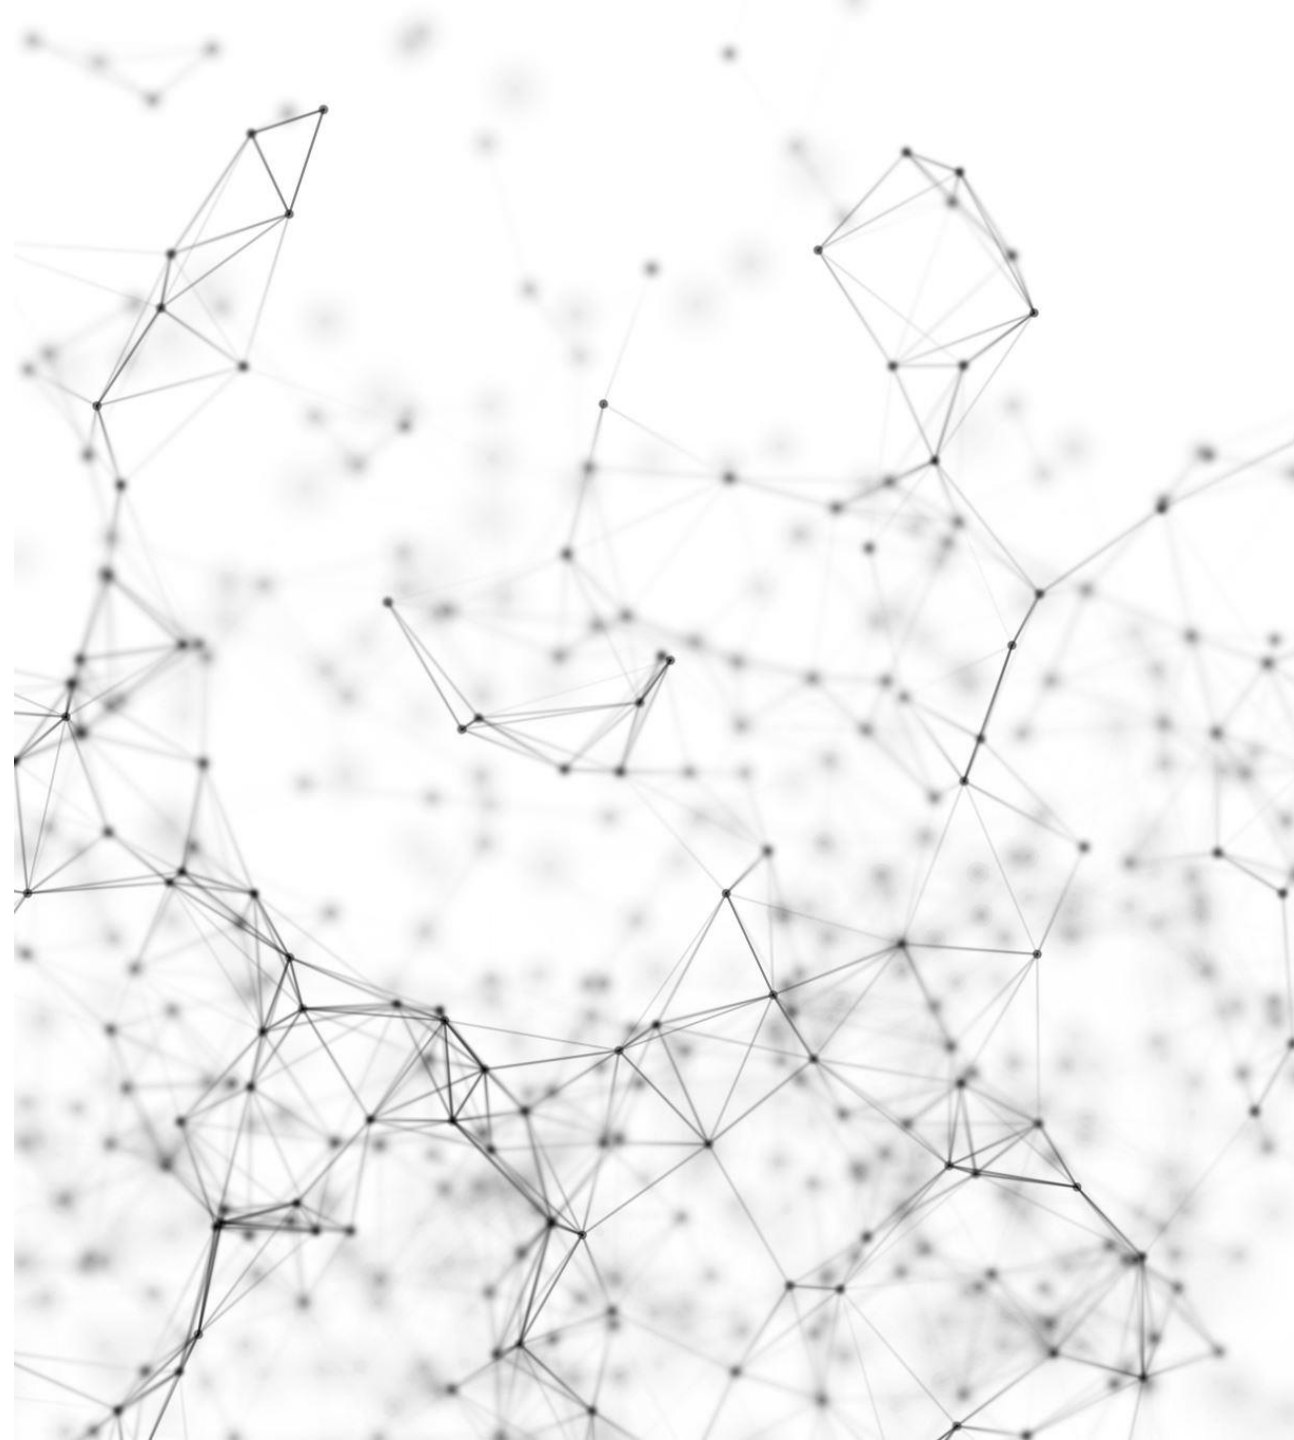


Multivariate Time Series

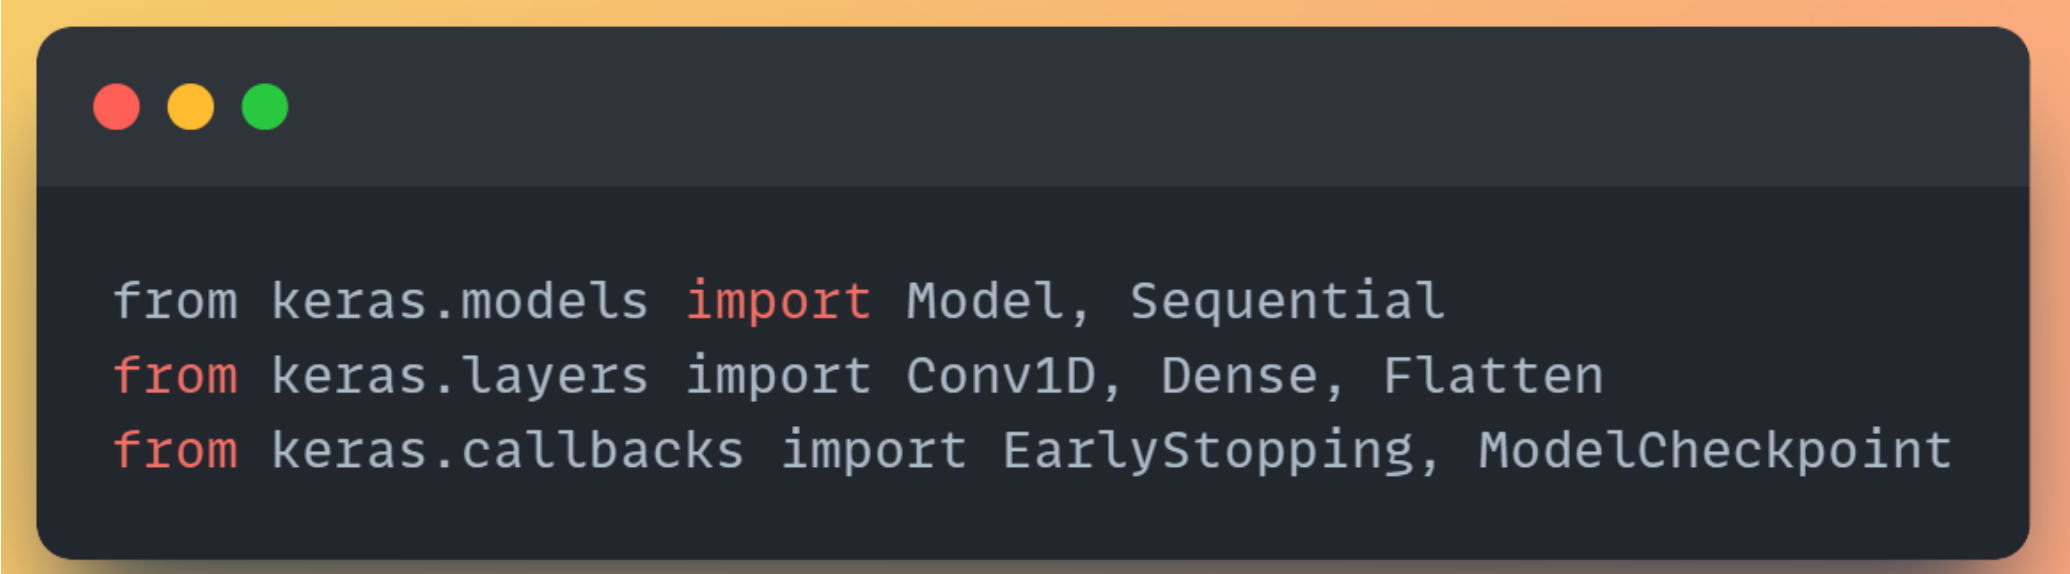
- Feature extraction with multiple filters (per dimension)



Convolutional neural networks (CNN) with Keras




Imports the KERAS modules (NN components)



```
from keras.models import Model, Sequential
from keras.layers import Conv1D, Dense, Flatten
from keras.callbacks import EarlyStopping, ModelCheckpoint
```

KERAS – 1D convolutional filter

TIME SERIES PAST VALUE WINDOW (T)




```
model = Sequential()
model.add(Conv1D(FILTERS_NUMBER, kernel_size=KERNEL_SIZE,
                 padding='causal', strides=1, activation='relu',
                 dilation_rate=1, input_shape=(T, 1)))
```

https://keras.io/api/layers/convolution_layers/convolution1d/

KERAS – 1D convolutional filter

TIME SERIES variable (1 = univariate case)



```
model = Sequential()
model.add(Conv1D(FILTERS_NUMBER, kernel_size=KERNEL_SIZE,
                 padding='causal', strides=1, activation='relu',
                 dilation_rate=1, input_shape=(T, 1)))
```

https://keras.io/api/layers/convolution_layers/convolution1d/

Do not forget...flatten + dense layer for the final prediction



```
model.add(Flatten())  
model.add(Dense(HORIZON, activation='linear'))
```



Number of predictions

https://keras.io/api/layers/reshaping_layers/flatten/

https://keras.io/api/layers/core_layers/dense/

Early stopping + best model save + training

```
model.compile(optimizer='Adam', loss='mse')

earlystop = EarlyStopping(monitor='val_loss', min_delta=0, patience=5)

best_val = ModelCheckpoint('model_{epoch:02d}.keras', save_best_only=True, mode='min')

history = model.fit(X_train,
                    y_train,
                    batch_size=BATCH_SIZE,
                    epochs=EPOCHS,
                    validation_data=(X_valid, y_valid),
                    callbacks=[earlystop, best_val],
                    verbose=1)
```

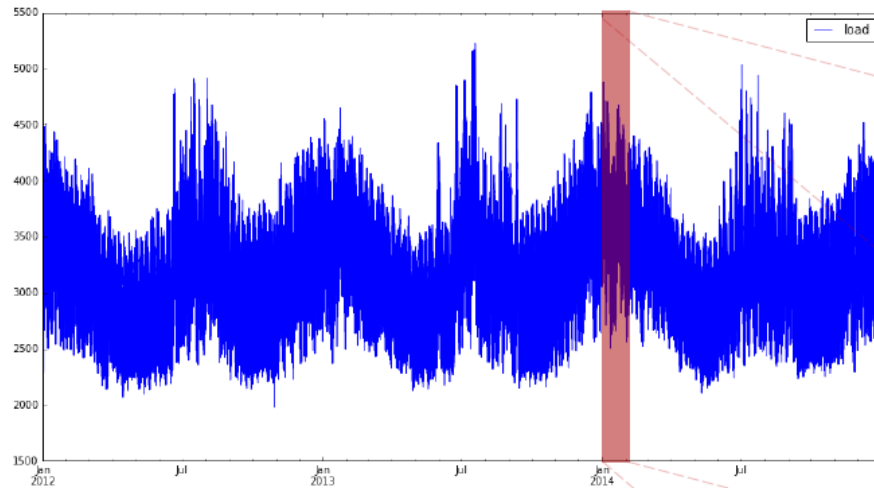
https://keras.io/api/models/model_training_apis/
https://keras.io/api/callbacks/early_stopping/

A black and white photograph of a wooden floor with white chalk markings. The markings include the letters 'FF' and the number '0'. The text 'Presentation du TD' is overlaid in the center.

Presentation du TD

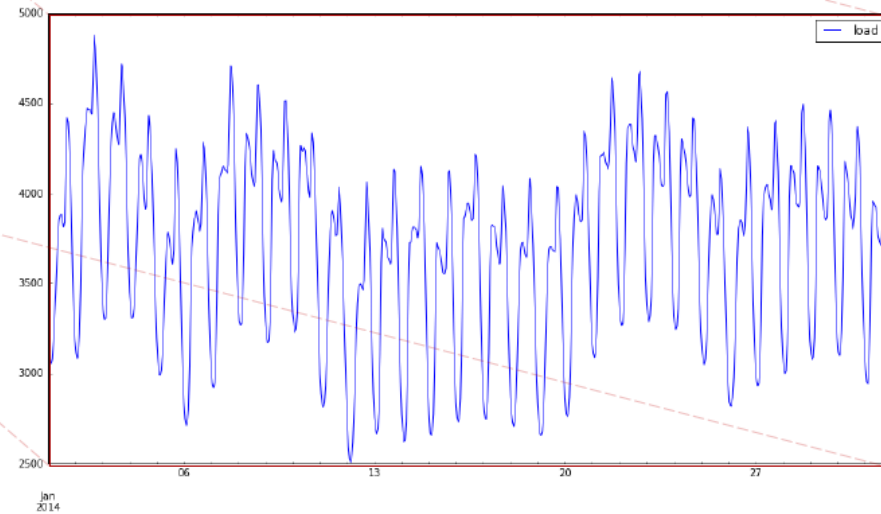
Data

New England ISO data



- 26,000 hourly load values
- Annual, weekly and daily seasonality

Tao Hong, Pierre Pinson, Shu Fan, Hamidreza Zareipour, Alberto Troccoli and Rob J. Hyndman, "Probabilistic energy forecasting: Global Energy Forecasting Competition 2014 and beyond", International Journal of Forecasting, vol.32, no.3, pp 896-913, July-September, 2016.



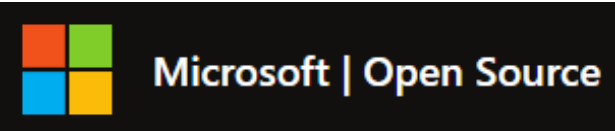
Notebook

- Open the file (Python Notebook):

TD_CNN_dilated_TS_FORECASTING.ipynb

- Instruction are contained in the notebook

References

- <https://stanford.edu/~shervine/teaching/cs-230/cheatsheet-deep-learning-tips-and-tricks>
- **DeepLearningForTimeSeriesForecasting (Microsoft)** 
<https://github.com/Azure/DeepLearningForTimeSeriesForecasting>
- **Ben Auffarth** Machine Learning for Time-Series with Python
- <https://github.com/PacktPublishing/Machine-Learning-for-Time-Series-with-Python/tree/main/chapter10>

