



# Time Series Management

Michele Linardi Ph.D.

[michele.linardi@orange.fr](mailto:michele.linardi@orange.fr)

Some of the slides of this course are taken from the **excellent** Tutorial of **Eammon Keogh**  
**A Decade of Progress in Indexing and Mining Large Time Series Databases**. VLDB 2006.



# Syllabus

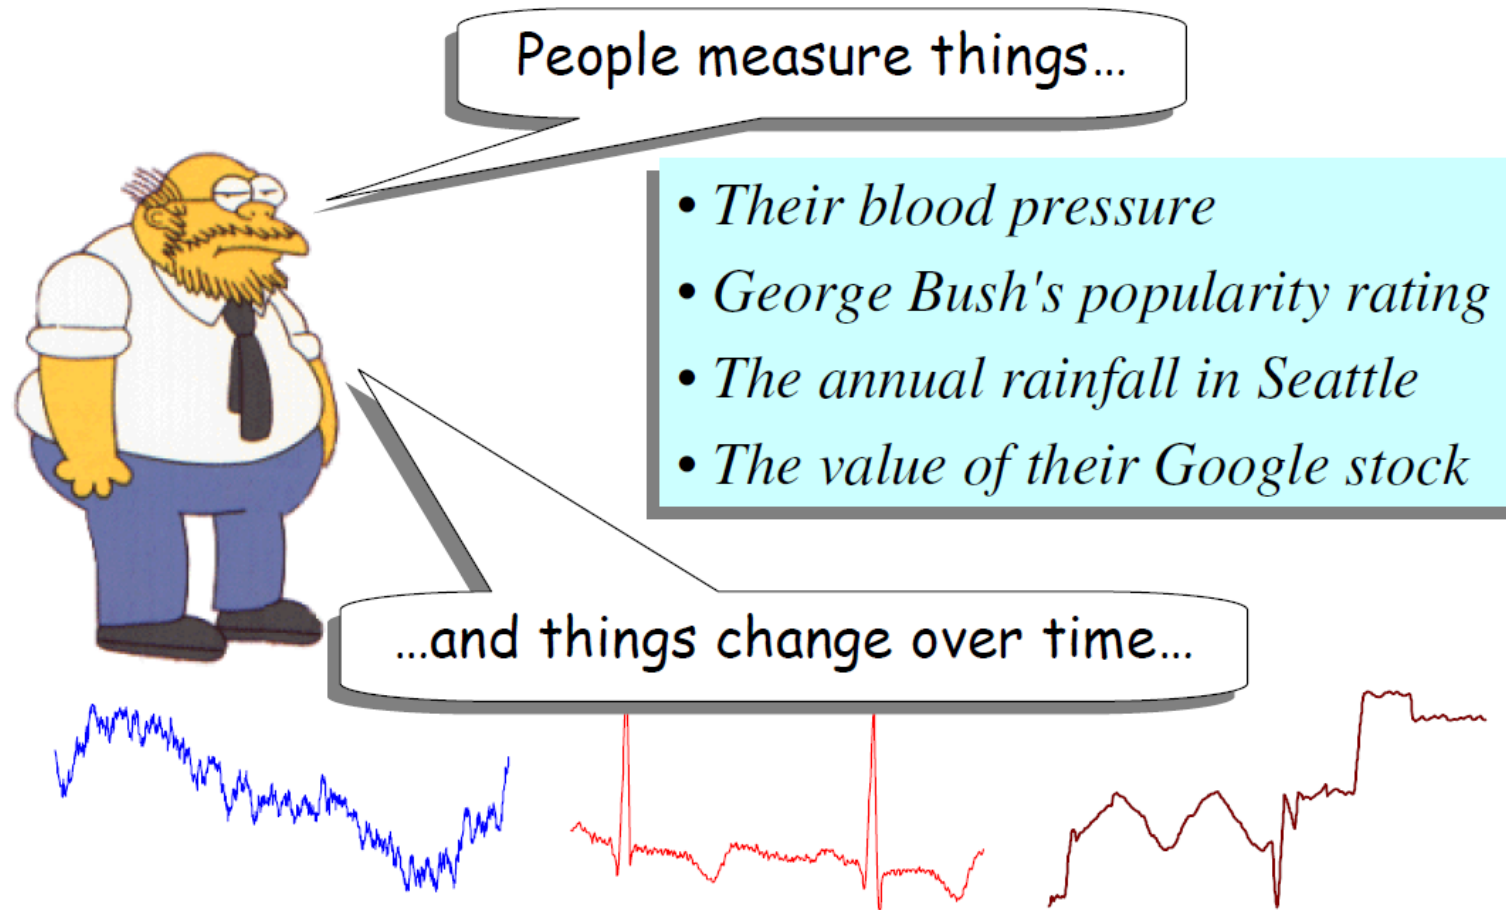
- Ubiquity of data series collections
- Time series data mining
- Similarity Search
- Metrics
- SAX : Time Series Symbolic Aggregate approXimation.
- DTW Lower bounding
- Speed up computation  
by lower bounding true Euclidean distance over SAX  
representation.

# Time series data... quick recap

- A **univariate time series** is a sequence of measurements of the same variable collected over time. Most often, the measurements are made at regular time intervals.



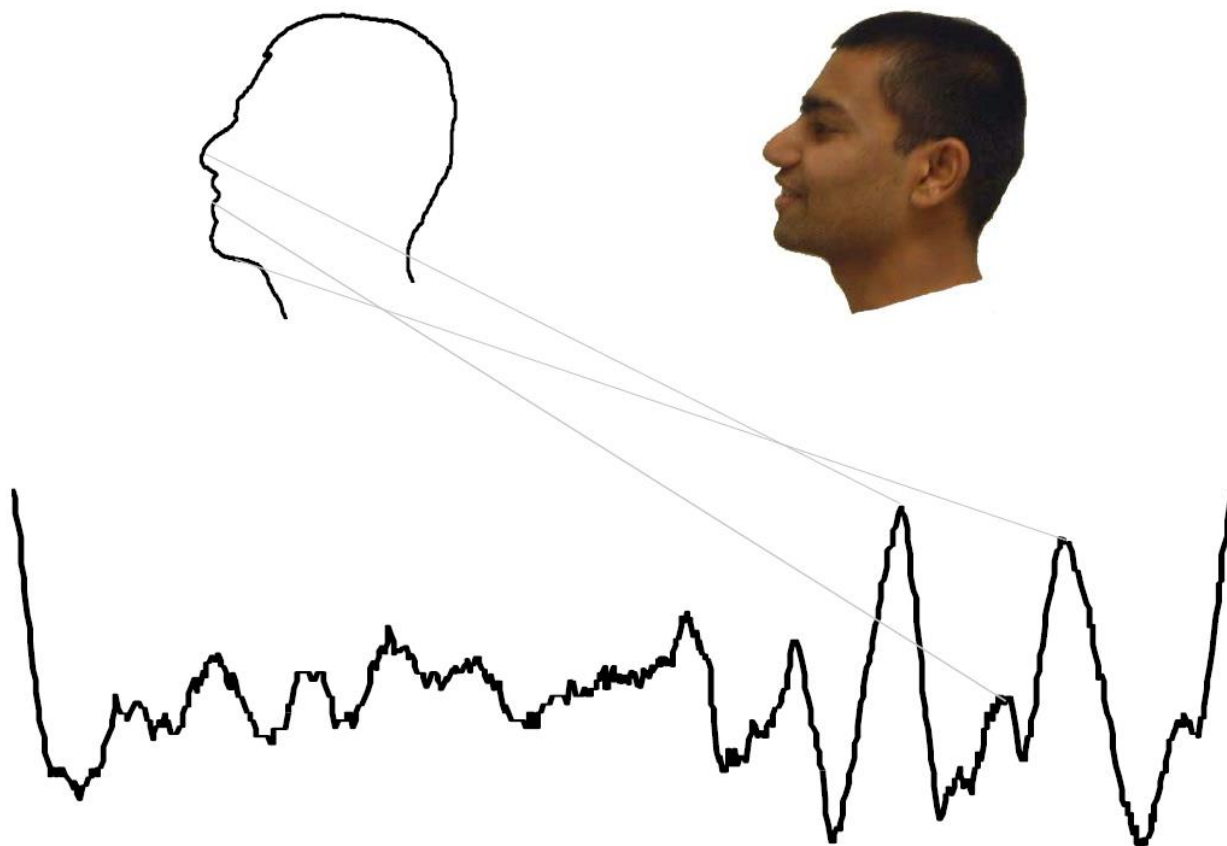
# Time series are ubiquitous



Thus time series occur in virtually every medical, scientific and businesses domain

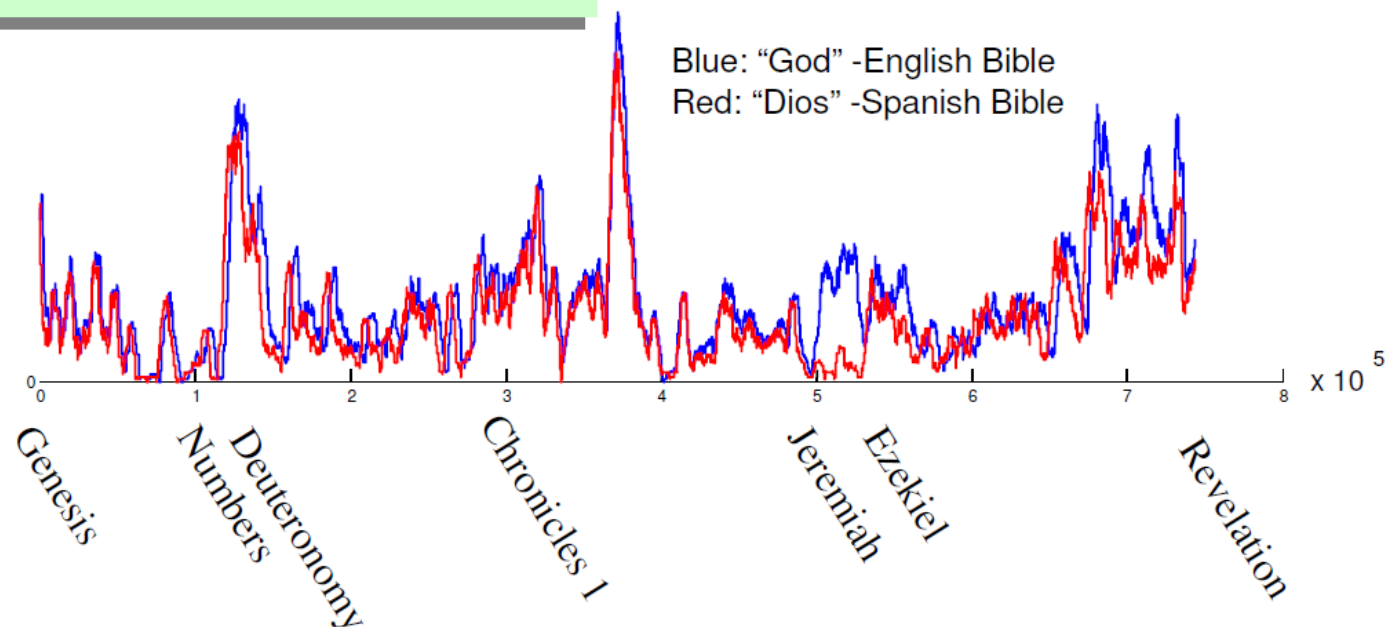
# Time series are ubiquitous (1/4)

Image data, may best be thought of as time series...



# Time series are ubiquitous (2/4)

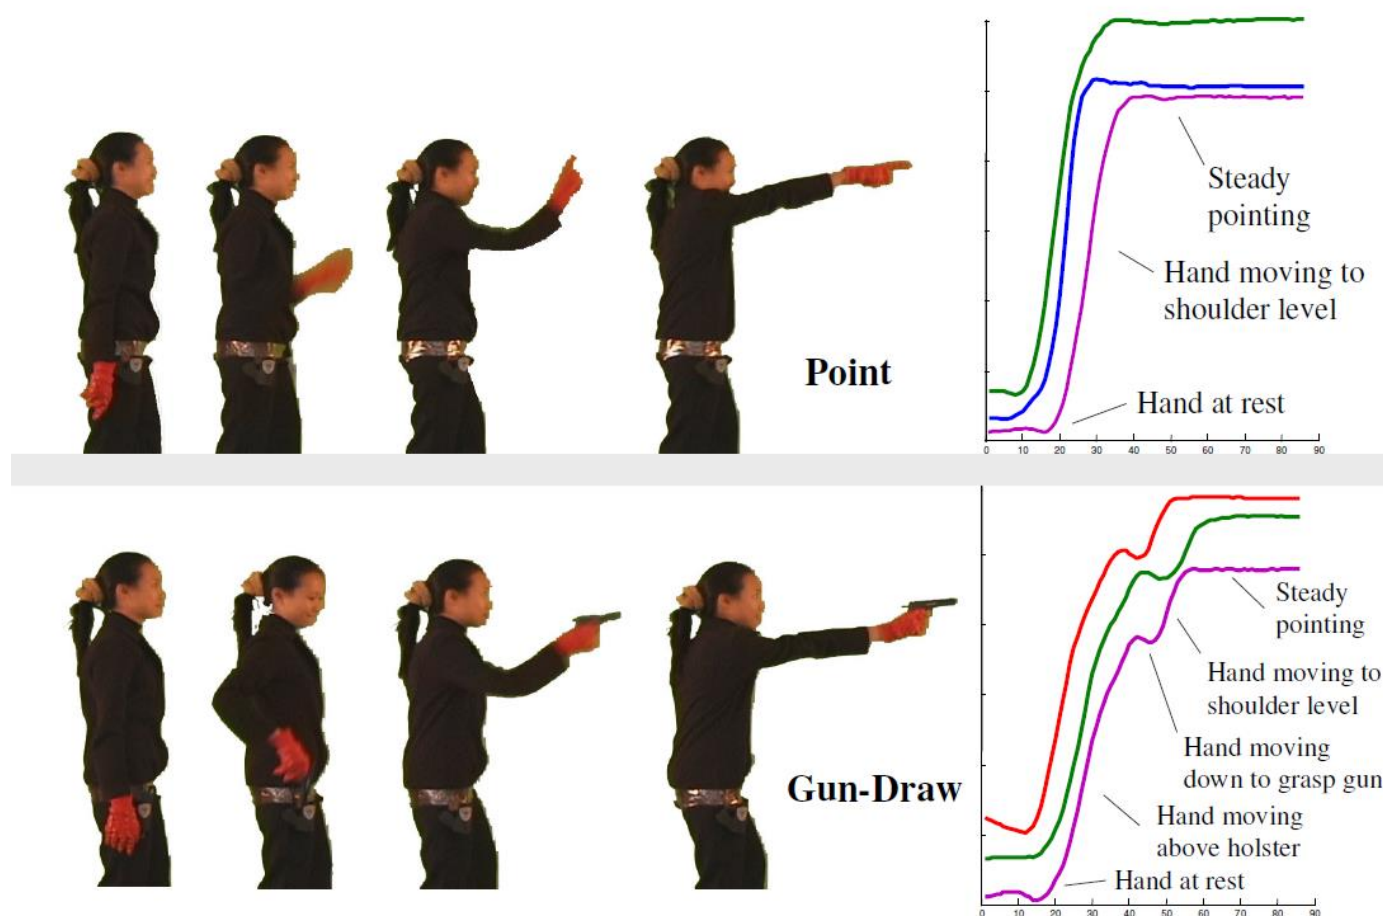
The local frequency  
of words in the Bible



Gray: "El Senor" -Spanish Bible

# Time series are ubiquitous (3/4)

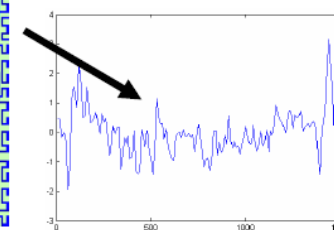
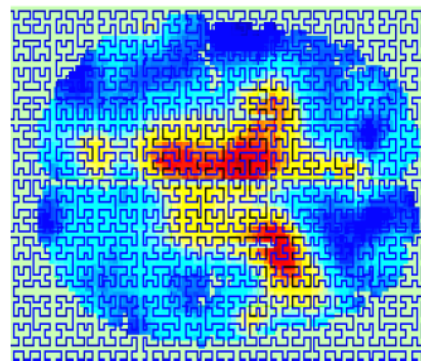
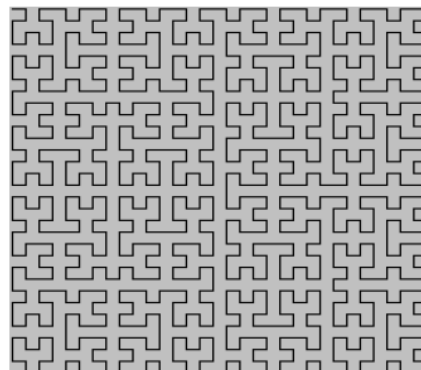
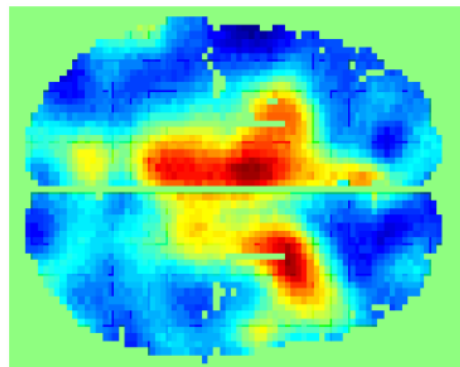
Video data, may best be thought of as time series...





# Time series are ubiquitous (4/4)

Brain scans (3D voxels), may best be thought of as time series..





# Why is Working With Time Series is so difficult? (1/3)

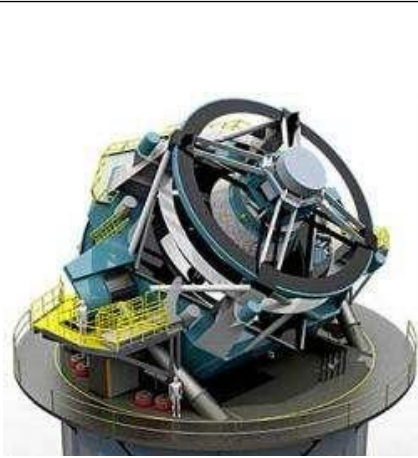
1 Hour of EKG data: 1 Gigabyte.

Typical Weblog: 5 Gigabytes per week.

Space Shuttle Database: 200 Gigabytes and growing.

Macho Database (Canadian Astronomy Data Centre): 3 Terabytes, updated with 3 gigabytes a day.

# Why is Working With Time Series is so difficult? (2/3)

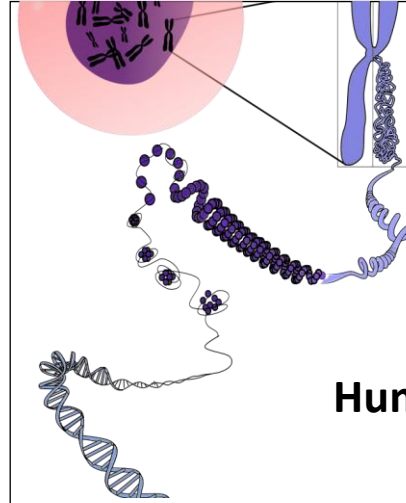


NASA's Solar Observatory

**1.5 TB per day**

Large Synoptic Survey  
Telescope (2019)

**~30 TB per night**



Human Genome project

**130 TB**



passenger aircrafts  
**20 TB per hour**

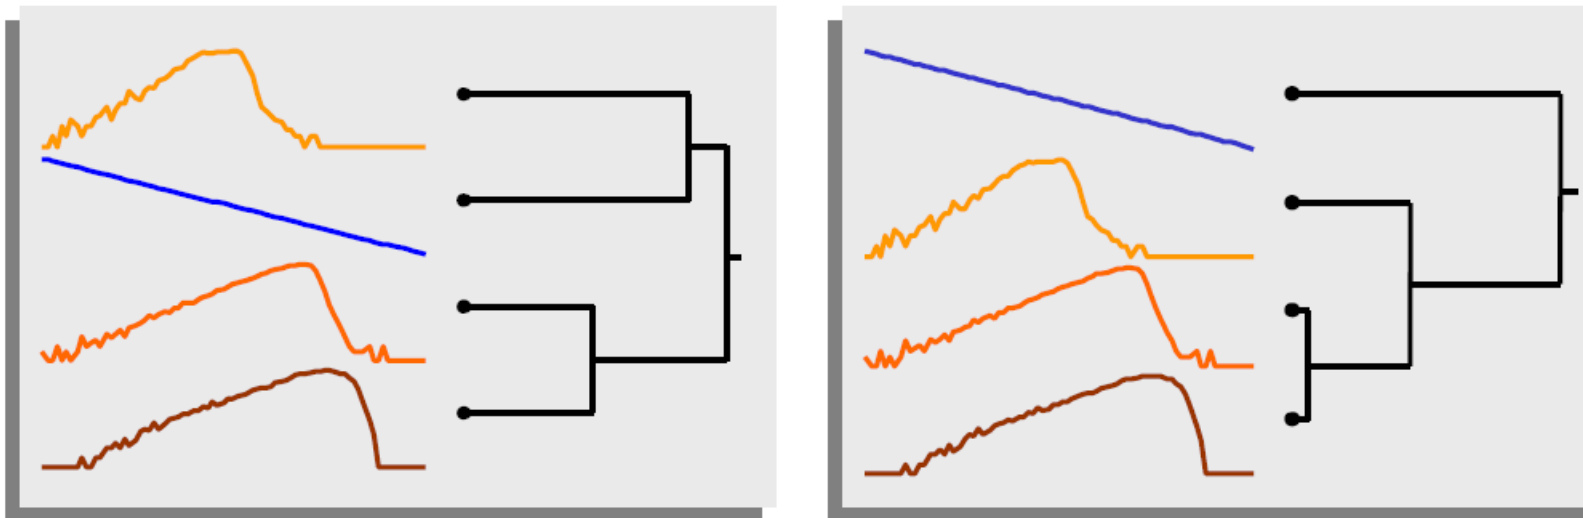
data center and  
services monitoring

**2B data series**  
**4M points/sec**



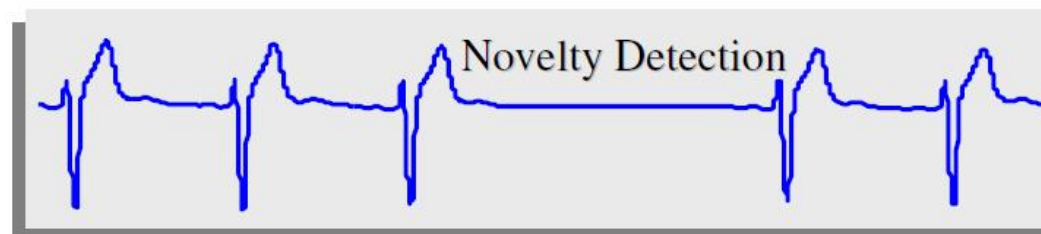
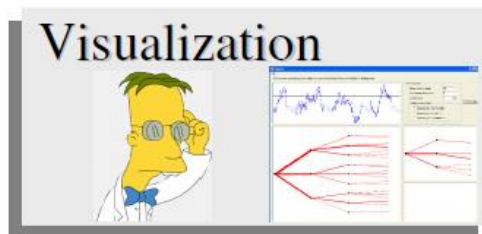
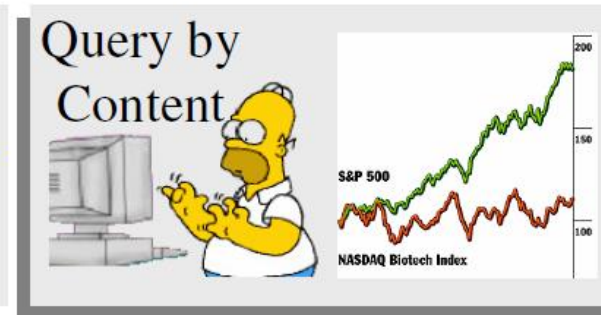
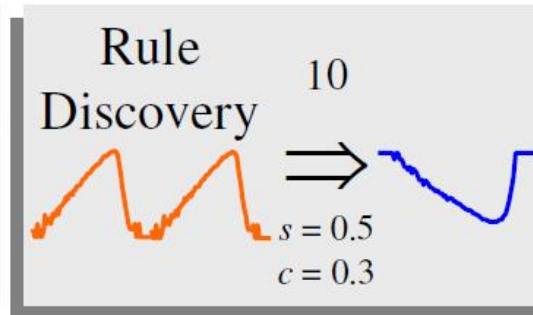
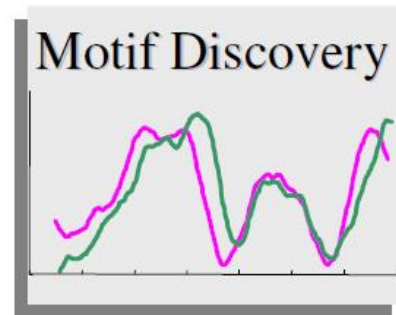
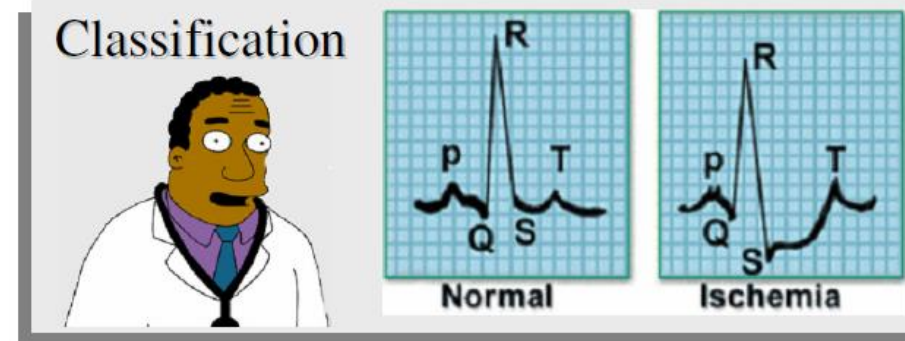
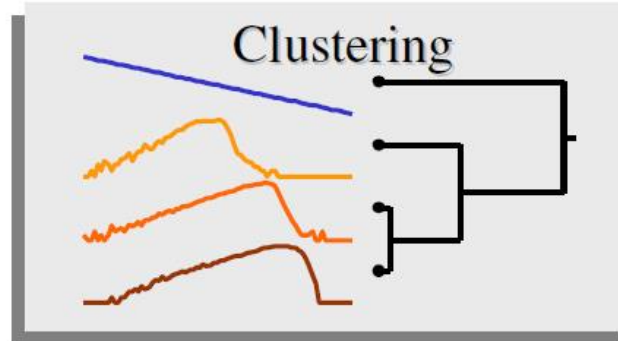
# Why is Working With Time Series is so difficult? (3/3)

**Answer:** We are dealing with subjectivity



The definition of similarity depends on the user, the domain and the task at hand. We need to be able to handle this subjectivity.

# Problems requiring Similarity Search



# Important Data Mining questions



How do we define similarity ?



How do we search large time series  
collection quickly ?

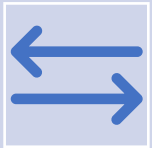
# What is similarity ?



# What is similarity ?



*The quality or state of being similar; likeness; resemblance; as, a similarity of features.*



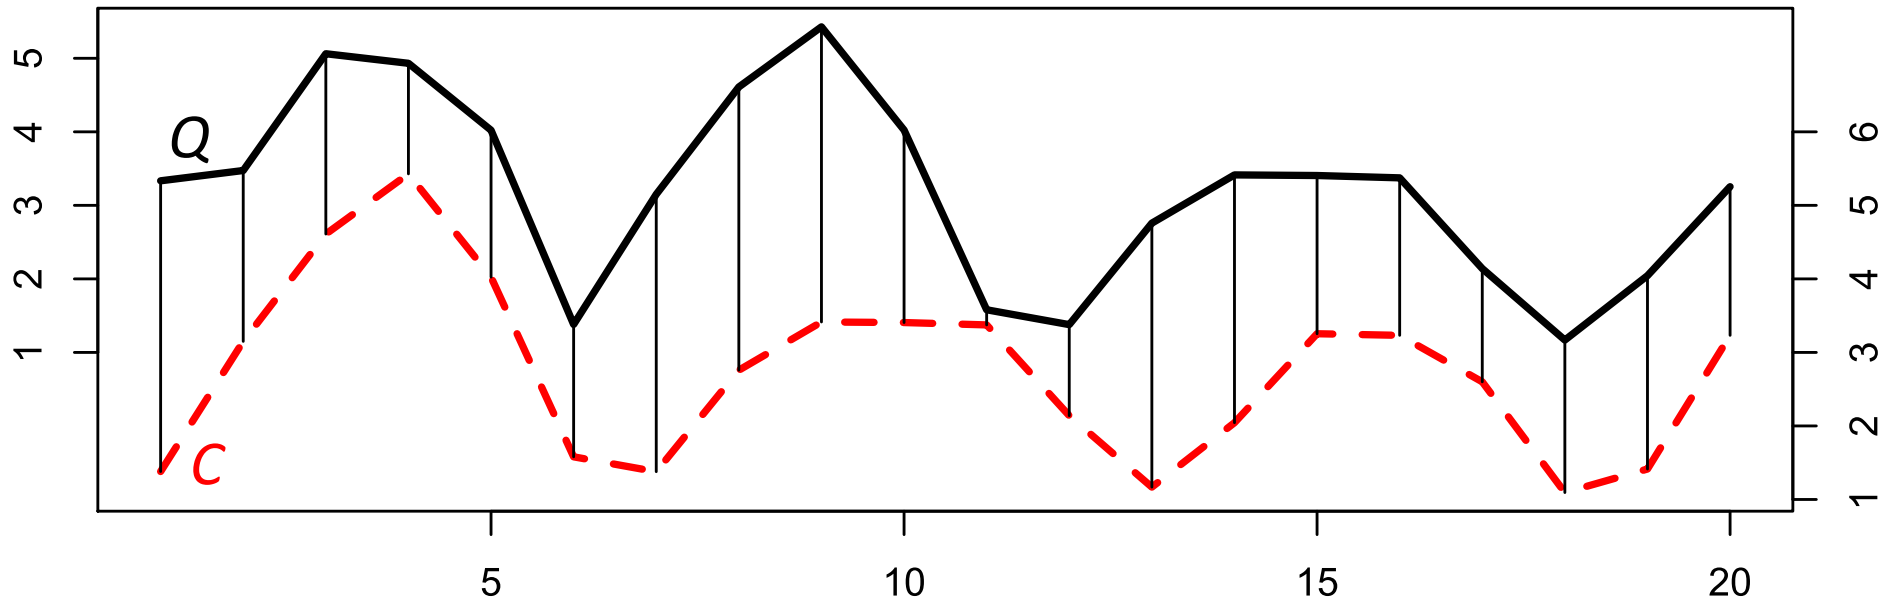
Similarity is hard to define, but...“We know it when we see it”



The real meaning of similarity is a philosophical question.



# Similarity Measure: Euclidean



Given two time series  $Q = q_1 \dots q_n$  and  $C = c_1 \dots c_n$

their Euclidean distance is defined as:

$$ED(Q, C) = \sqrt{\sum_{i=1}^n (q_i - c_i)^2}$$

# Preprocessing the data before distance calculations

- Z-Normalization (Amplitude Scaling)
- Linear Trend
- Noise



# Z Normalization

Z-normalization, also known as

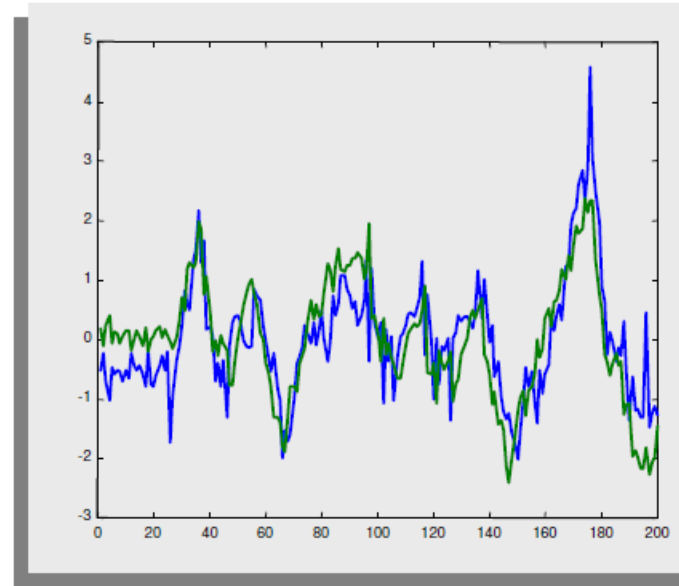
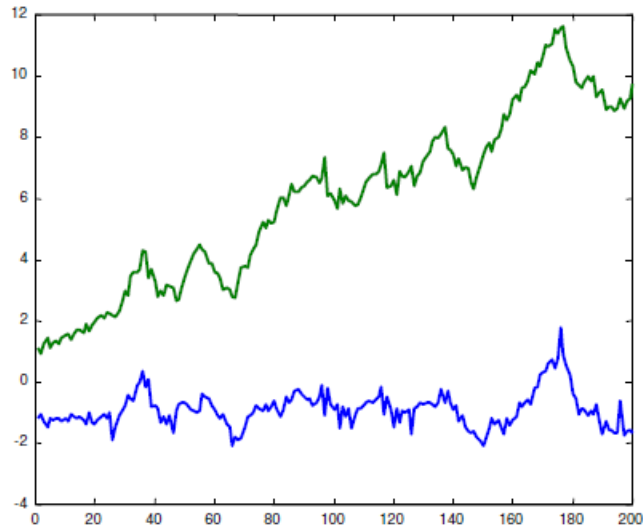
**“Normalization to Zero Mean and Unit of Energy”.**

The procedure ensures, that all elements of the input vector are transformed into the output vector whose mean is approximately 0 while the standard deviation is in a range close to 1.

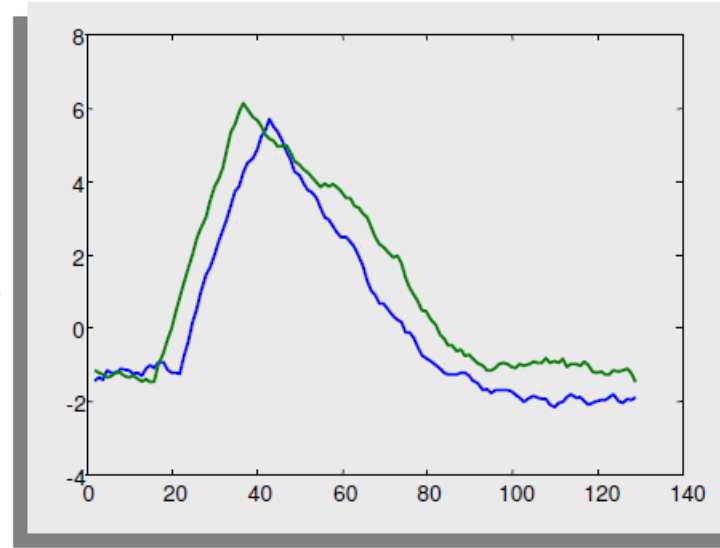
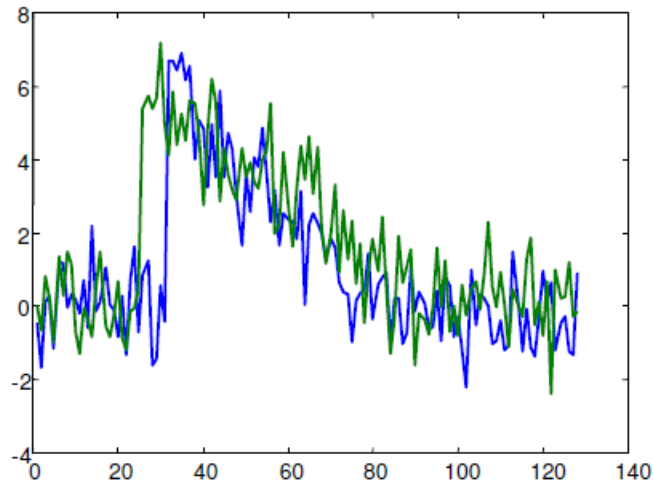
$$x'_i = \frac{x_i - \mu}{\sigma}, \text{ where } i \in \mathbb{N}$$



# Remove linear trend



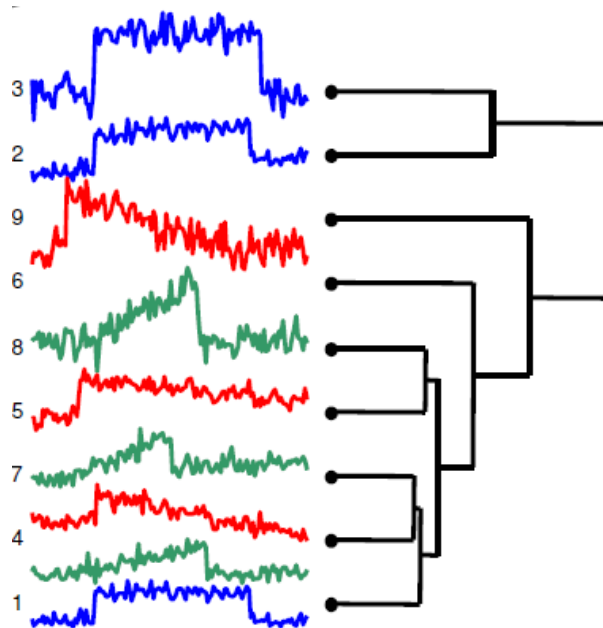
# Removing Noise




Smoothing function: Remove noise component  
e.g., Average each datapoints value with its neighbors

# Importance of data pre-processing



**Clustered using Euclidean distance on the raw data.**



**Clustered using Euclidean distance, after removing noise, linear trend, and Z-Normalization**

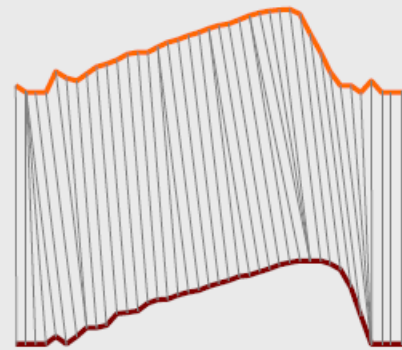


# Dynamic Time Warping



Fixed Time Axis

*Sequences are aligned "one to one".*





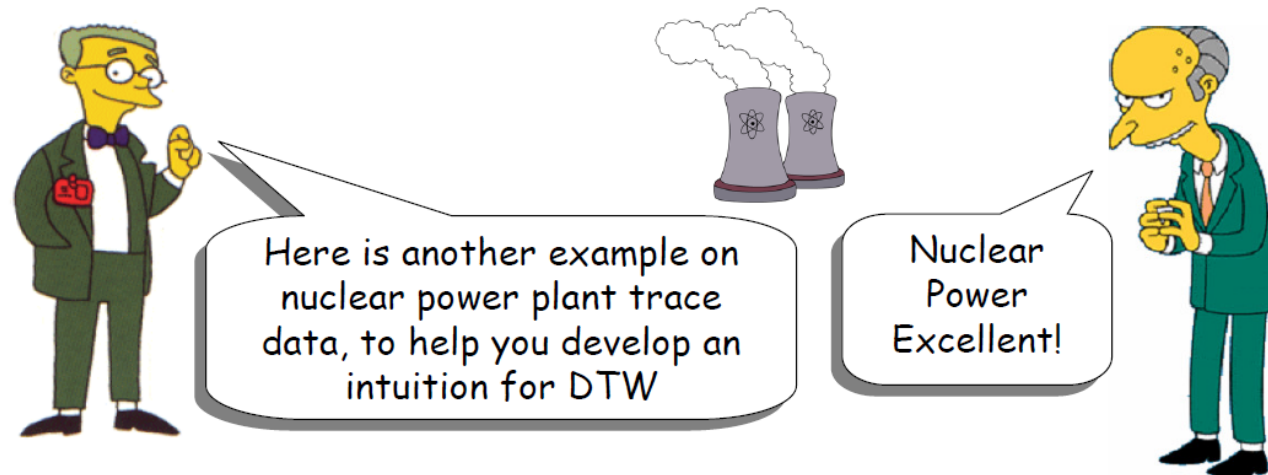
"Warped" Time Axis

*Nonlinear alignments are possible.*

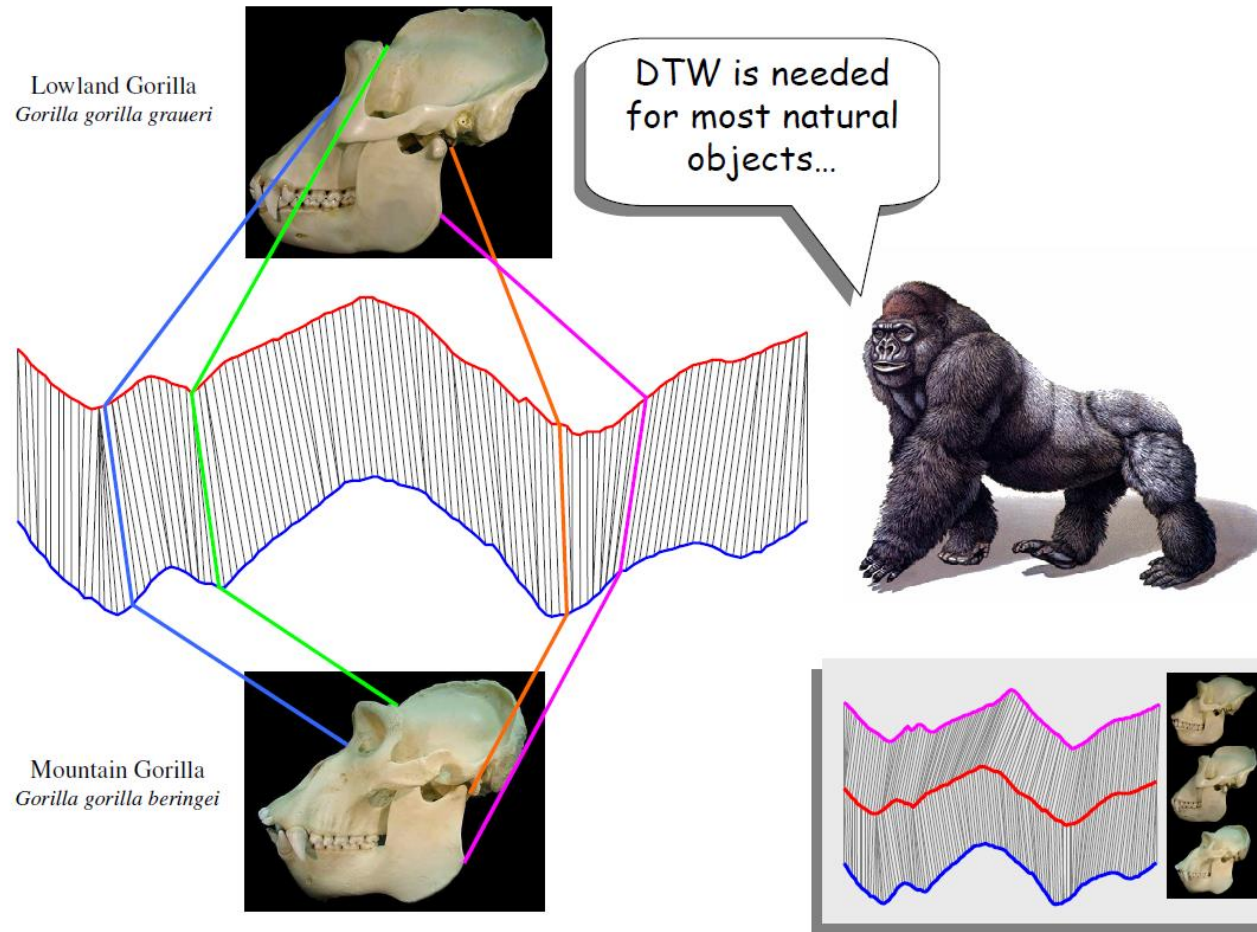
Note: We will first see the utility of DTW, then see how it is calculated.




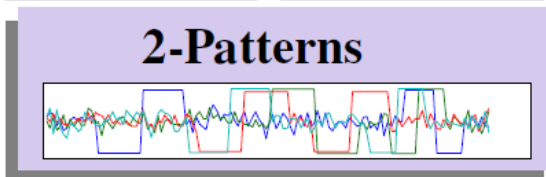



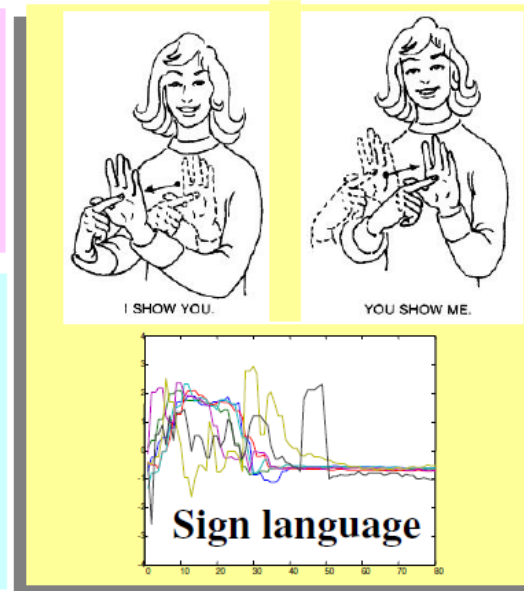
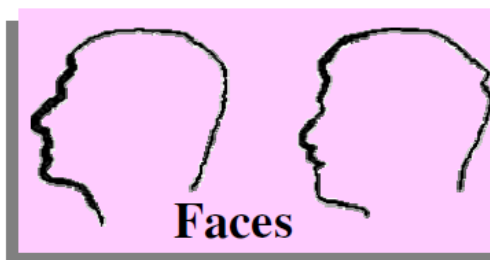
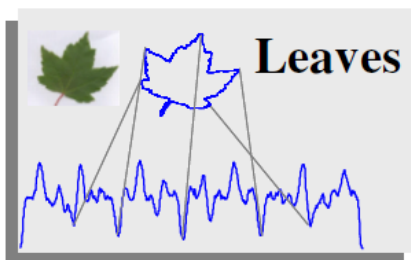
# Some real-world example (1/3)




# Some real-world example (2/3)



# Some real-world example (3/3)



# Distance Measure: Dynamic Time Warping (DTW)




Given two time series  $Q = q_1 \dots q_n$  and  $C = c_1 \dots c_n$

We want to find a pairwise points alignment  $w \in (\mathbb{N} \times \mathbb{N})^n$ , which minimize the pairwise points distance:

$$\text{DTW}(Q, C) = \underset{P}{\operatorname{argmin}} \left( \sqrt{\sum_{i=1}^{|P|} (q_{P_i[0]} - c_{P_i[1]})^2} \right)$$

# DTW calculation




Recursive function of minimum cost path

$$\gamma(P_i) = \sqrt{(q_{P_i[0]} - c_{P_i[1]})^2 + \min\{\gamma(P_{i[0]} - 1), \gamma(P_{i[1]} - 1), \gamma(P_{i[0]} - 1, P_{i[1]} - 1)\}}$$

**$P_0 = (0, 0)$  and  $P_n = (n, n)$  are always the first and the last element of the warping path respectively.**

# Time complexity – $O(n^2)$




- Dynamic Time Warping gives **much better** results than Euclidean distance on virtually all problems.
- Dynamic Time Warping is very very slow to calculate!

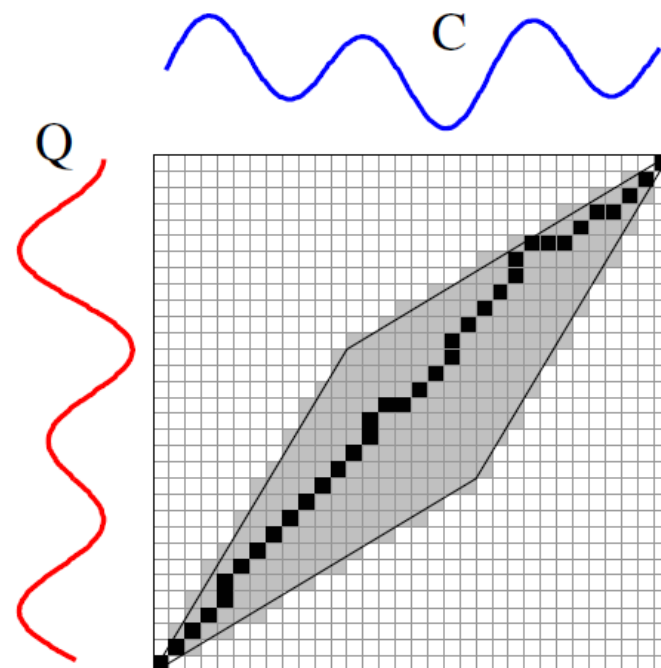
Is there anything we can do to speed up similarity search under DTW?

# Global warping path constraints

- Slightly speed up the calculations
- Prevent pathological warpings



Sakoe-Chiba Band



Itakura Parallelogram




# Lower bounding


Assume that we have two functions:

- $\text{DTW}(A,B)$
- $\text{lower\_bound\_distance}(A,B)$

The true DTW  
function is very  
slow...



The *lower  
bound* function  
is very fast...



By definition, for all  $A, B$ , we have

$$\text{lower\_bound\_distance}(A,B) \leq \text{DTW}(A,B)$$

# Speed up search with lower bounding

We can speed up similarity search under DTW  
by using a lower bounding function

## Algorithm Lower\_Bounding\_Sequential\_Scan(Q)

```
1.  best_so_far = infinity;  
2.  for all sequences in database  
3.    LB_dist = lower_bound_distance(Ci, Q);  
4.    if LB_dist < best_so_far  
5.      true_dist = DTW(Ci, Q);  
6.      if true_dist < best_so_far  
7.        best_so_far = true_dist;  
8.        index_of_best_match = i;  
9.      endif  
10.   endif  
11. endfor
```

Try to use a cheap  
lower bounding  
calculation as  
often as possible.

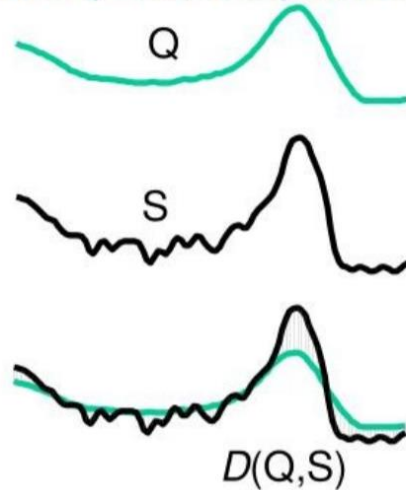


Only do the  
expensive, full  
calculations when  
it is absolutely  
necessary



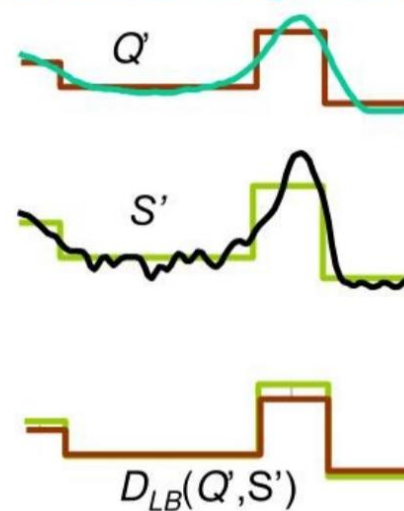
# Lower Bounding Euclidean distance

Exact (Euclidean) distance  $D(Q,S)$



$$D(Q,S) \equiv \sqrt{\sum_{i=1}^n (q_i - s_i)^2}$$

Lower bounding distance  $D_{LB}(Q,S)$

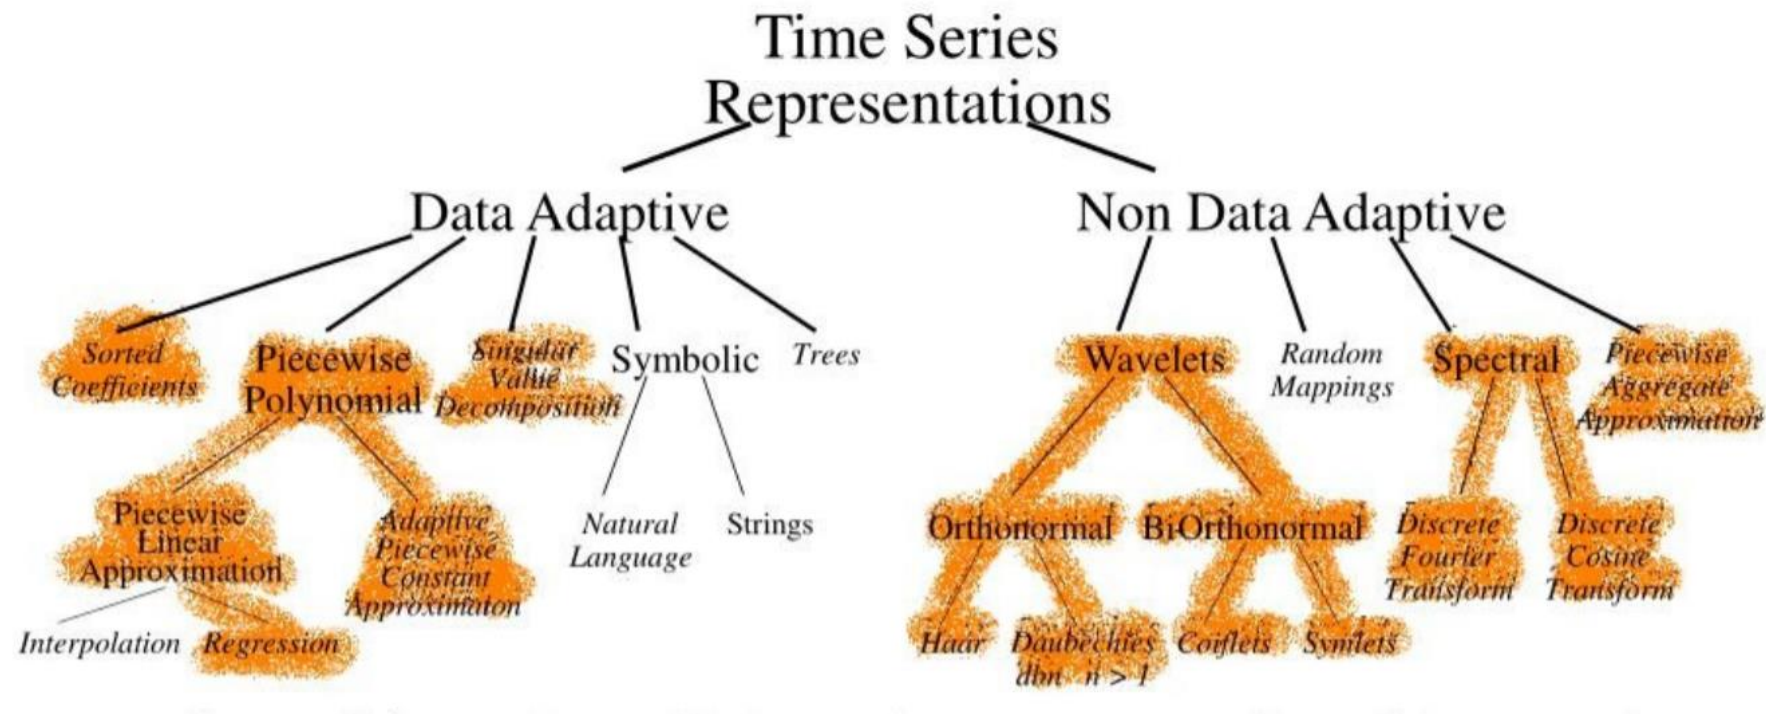


$$D_{LB}(Q',S') \equiv \sqrt{\sum_{i=1}^M (sr_i - sr_{i-1})(qv_i - sv_i)^2}$$

Lower bounding means that for all Q and S, we have...

$$D_{LB}(Q',S') \leq D(Q,S)$$

# Lower bounding over discrete TS representation



# Symbolic Aggregate approXimation (SAX)




baabccbc

# Symbolic Aggregate approXimation (SAX)




How do we obtain SAX?



First convert the time series to PAA representation, then convert the PAA to symbols

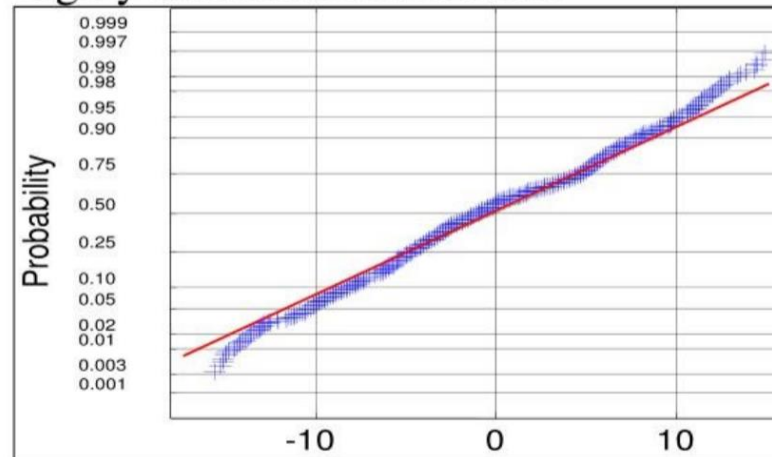
It take linear time



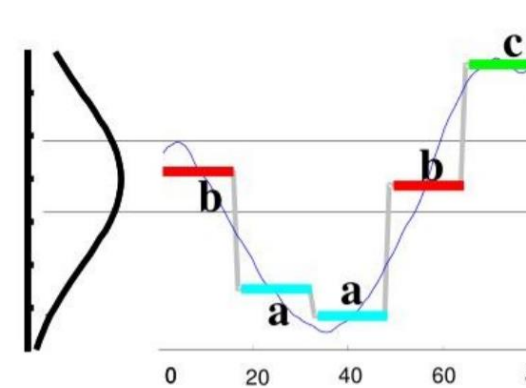
baabccbc

# Symbolic Aggregate approXimation (SAX)

Time series subsequences tend to have a highly Gaussian distribution



A normal probability plot of the (cumulative) distribution of values from subsequences of length 128.

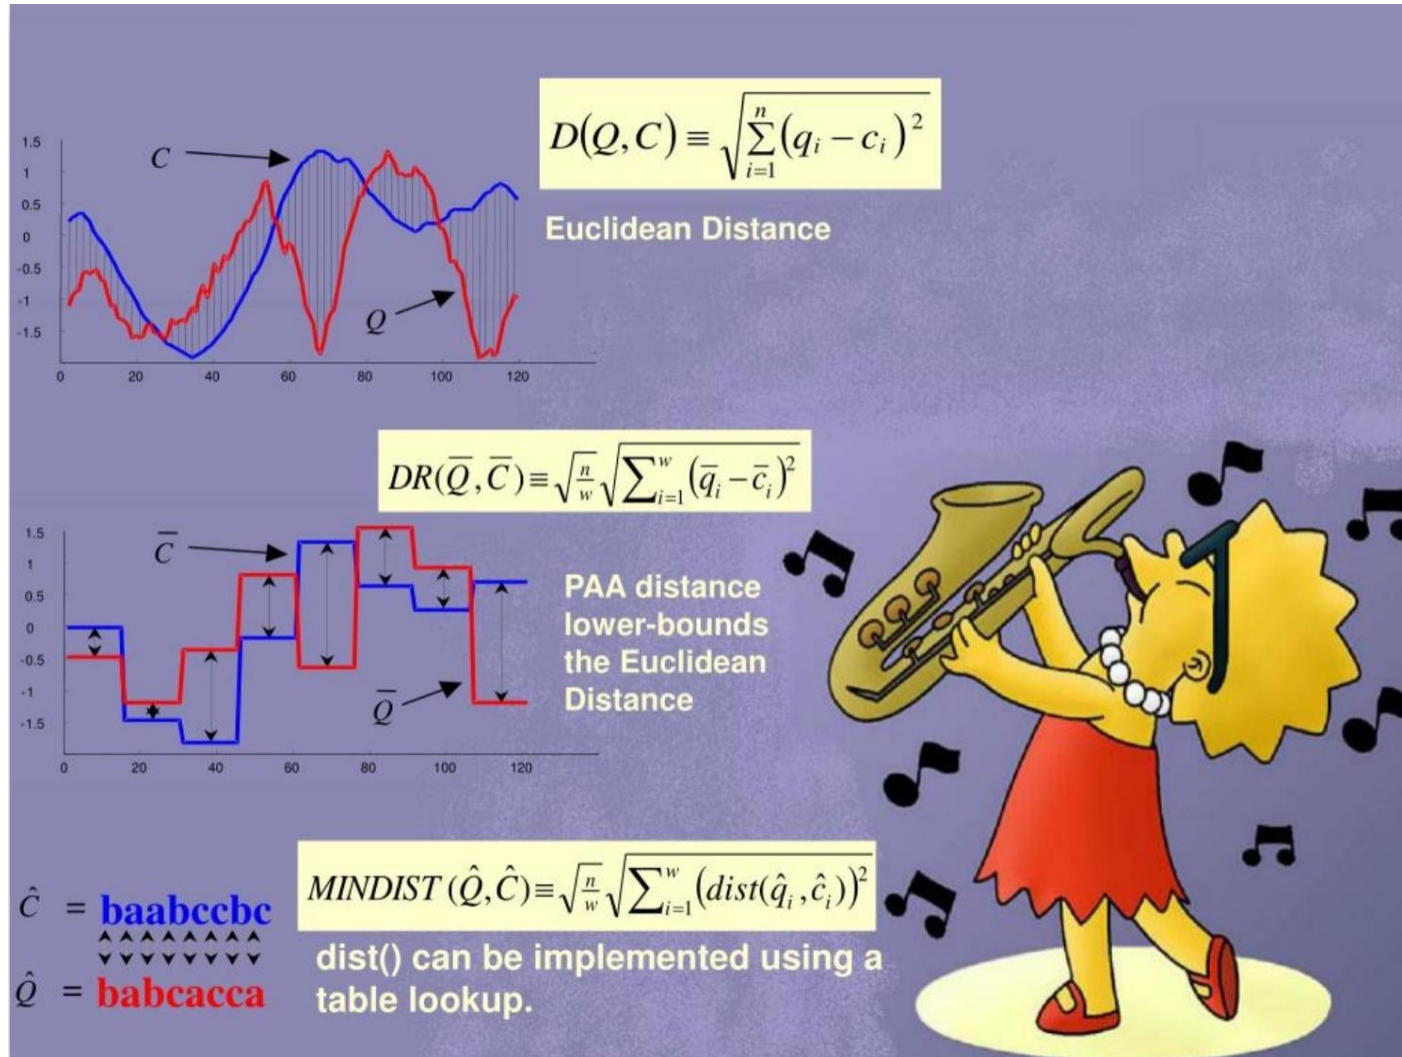


Why a Gaussian?





# SAX lowerbounding (MINDIST)



# Implementation (PYTS library)



A Python Package for Time Series Classification

Star 1,615

## Navigation

### Getting Started

[Installation, testing and development](#)

[Contributing guide](#)

### Documentation

[User guide](#)

[API Documentation](#)

[Scikit-learn compatibility](#)

### Tutorial - Examples

[Introductory examples](#)

[Approximating time series](#)

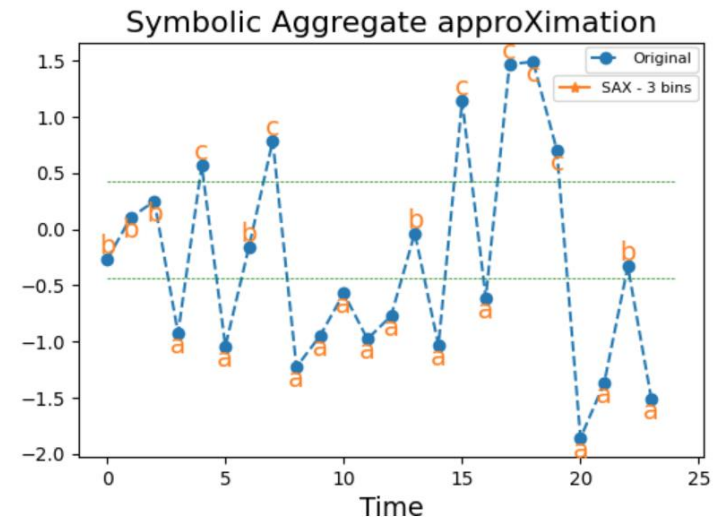
[Bag-of-words transformation](#)

[Classification algorithms](#)


**Note:** Click [here](#) to download the full example code

## Symbolic Aggregate approXimation

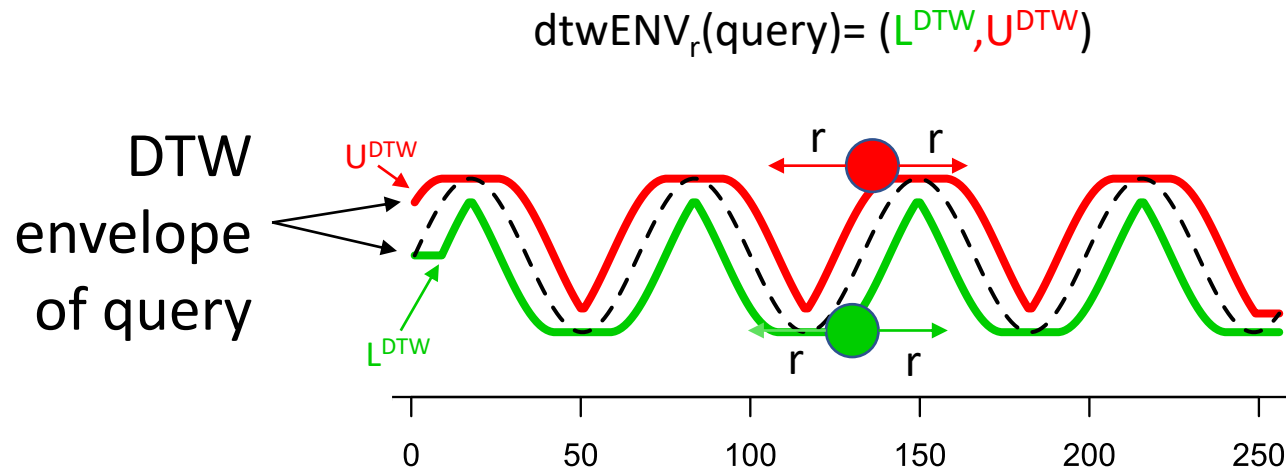
Binning continuous data into intervals can be seen as an approximation that reduces noise and captures the trend of a time series. The Symbolic Aggregate approXimation (SAX) algorithm bins continuous time series into intervals, transforming independently each time series (a sequence of floats) into a sequence of symbols, usually letters. This example illustrates the transformation. It is implemented as `pyts.approximation.SymbolicAggregateApproximation`.



# Lower bounding DTW




# Lower bounding DTW



- For each query position,  $\text{dtwENV}$  bounds the values that can be aligned to a single point in the data series candidates (**warping window  $r$** ).

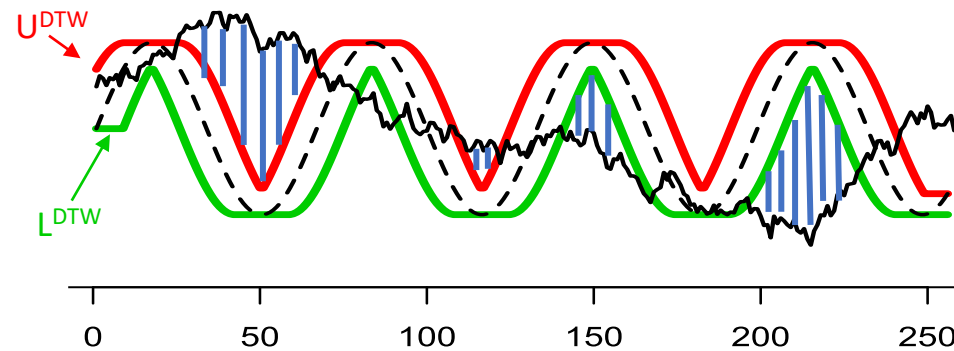
# Lower bounding DTW

$\text{LB\_Keogh}(\text{dtwENV}_r(\text{query}), C)$



# Lower bounding DTW

$\text{LB\_Keogh}(\text{dtwENV}_r(\text{query}), \text{candidate})$



- Lower bounding of the true DTW distance between Query and Candidate  
[  $O(n)$  time ]

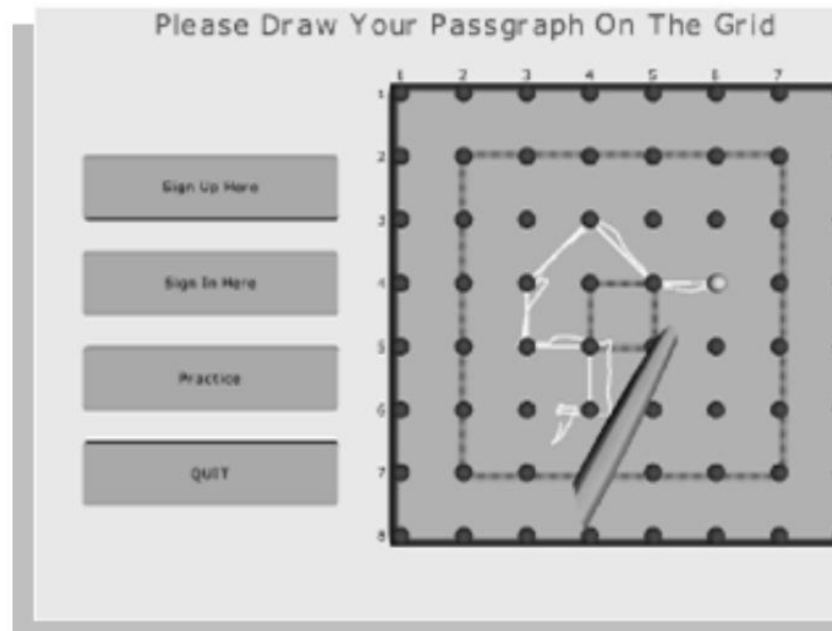


A black and white photograph of a wooden floor with white chalk markings. The markings include the letters 'FF' and the number '0'. The text 'Presentation du TD' is overlaid in the center.

# Presentation du TD

# Haptics DATA

Data are taken from 5 people entering their passgraph (a code to access a system protected by a graphical authentication system) on a touchscreen. The data are the x-axis movement only.





# Notebook

- Open the file (Python Notebook):

**TS\_SimilaritySearch.ipynb**

- Instruction are contained in the notebook
- You can look at the data source to find similarity search hyperparameter :

[https://www.cs.ucr.edu/%7Eeamonn/time\\_series\\_data\\_2018/](https://www.cs.ucr.edu/%7Eeamonn/time_series_data_2018/)

# References

- Eamonn J. Keogh. A Decade of Progress in Indexing and Mining Large Time Series Databases. VLDB 2006: 1268
- Eamonn J. Keogh, Li Wei, Xiaopeng Xi, Sang-Hee Lee, Michail Vlachos: **LB\_Keogh Supports Exact Indexing of Shapes under Rotation Invariance with Arbitrary Representations and Distance Measures. VLDB 2006: 882-893**
- Chotirat Ann Ralanamahatana, Jessica Lin, Dimitrios Gunopulos, Eamonn Keogh, Michail Vlachos & Gautam Das **Mining Time Series Data**  
[https://link.springer.com/chapter/10.1007/0-387-25465-x\\_51#Abs1](https://link.springer.com/chapter/10.1007/0-387-25465-x_51#Abs1)
- <https://tslearn.readthedocs.io/> (tslearn package documentation)