



Time Series Management

Michele Linardi Ph.D.

michele.linardi@orange.fr




Syllabus

- Time series data mining
- Motif primitive
- Discord (Outlier)
- Matrix profile Algorithm
- Algorithm Python code DEMO


Data Series Motifs

- **Motif discovery** is one of the most useful primitives for data series mining
- Given a time series T , a motif is the pair of subsequences of length ℓ with the smallest Euclidean distance.




Data Series Discords

- **Discords** are used to detect anomalous/interesting patterns
 - usually represent the most unusual subsequences within a **time series**
- Given a time series T , a discord is the subsequence of length ℓ that has the largest Euclidean distance to its nearest neighbor.



Matrix Profile - Motif and Discord Discovery




Distance Matrix: a symmetric matrix, which contains the pairwise subsequences (of length ℓ) distances in T .

Excluding Trivial matches : = distances of any sequence pair having indexes (e.g., i and j) closer than a given threshold

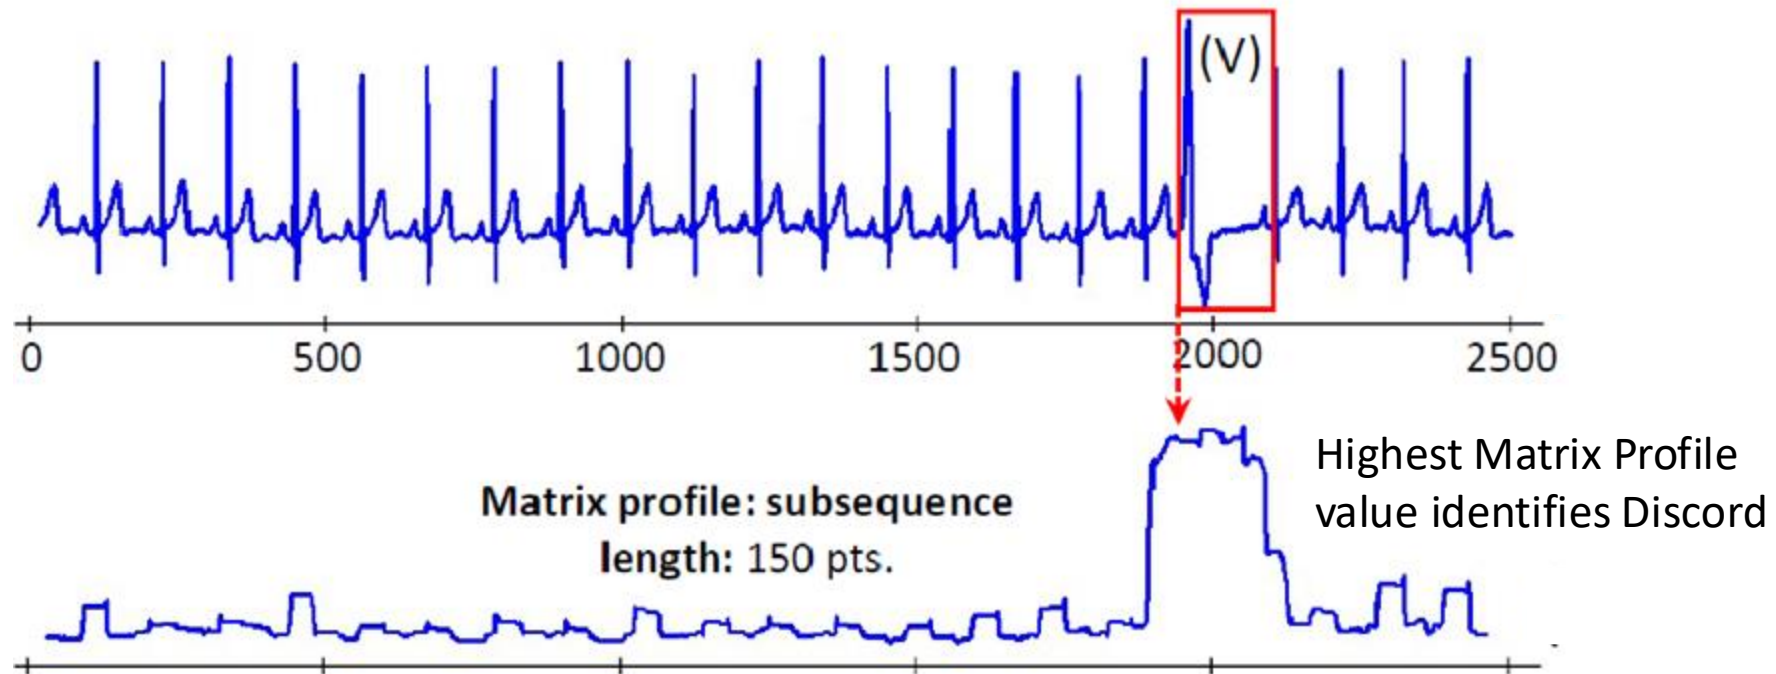
Matrix Profile: a vector of distance between each subsequence and its nearest neighbor

Matrix Profile - Motif and Discord Discovery



A pair of minimum points identifies the Motif

Matrix Profile - Motif and Discord Discovery



Euclidean distance metrics

Given two time series

$$\mathbf{x} = x_1 \dots x_n$$

and

$$\mathbf{y} = y_1 \dots y_n$$



their z-Normalized Euclidean distance is defined as:

$$\hat{x}_i = \frac{x_i - \mu_x}{\sigma_x} \quad \hat{y}_i = \frac{y_i - \mu_y}{\sigma_y}$$

```
function y = zNorm(x)
y = (x-mean(x))/std(x,1);
```

$$d(x,y) = \sqrt{\sum_{i=1}^n (\hat{x}_i - \hat{y}_i)^2}$$

```
function d = EuclideanDistance(x,y)
d = sqrt(sum((x-y).^2));
```



Pearson correlation coefficient

- Given two time series x and y of length m .
- Correlation Coefficient:

$$\text{corr}(x, y) = \frac{(E(x) - \mu_x)(E(y) - \mu_y)}{\sigma_x \sigma_y} = \frac{\sum_{i=1}^m x_i y_i - m \mu_x \mu_y}{m \sigma_x \sigma_y}$$

- Where $\mu_x = \frac{\sum_{i=1}^m x_i}{m}$ and $\sigma_x^2 = \frac{\sum_{i=1}^m x_i^2}{m} - \mu_x^2$

- Sufficient Statistics:

$$\sum_{i=1}^m x_i y_i \quad \sum_{i=1}^m x_i \quad \sum_{i=1}^m y_i \quad \sum_{i=1}^m x_i^2 \quad \sum_{i=1}^m y_i^2$$

The sufficient statistics can be calculated in one linear scan. Given the sufficient statistics, correlation coefficient is a constant operation. Note the use of the dot product, which is the key component of many lockstep measures.


Correlation – Euclidean Distance relationship

Z-normalized Euclidean distance

m := number of data points

$$d(\hat{x}, \hat{y}) = \sqrt{2m(1 - \text{corr}(x, y))}$$


- Correlation coefficient does not obey triangular inequality, while Euclidean distance does
- Maximizing correlation coefficient can be achieved by minimizing normalized Euclidean distance and vice versa
- Correlation coefficient is bounded between -1 and 1, while z-normalized Euclidean distance is bounded between zero and a positive number dependent on m



20 for $m = 100$

Mueen's Algorithm for Similarity Search (MASS)

- MASS uses a convolution based method to calculate sliding dot products in $O(n \log n)$
- Convolution: If x and y are vectors of polynomial coefficients, convolving them is equivalent to multiplying the two polynomials.
- We can use convolution to compute all of the sliding dot products between the query and sliding windows.
- Convolution can be computed in the frequency domain, using the **Fast Fourier Transform**.



Matrix profile algorithm – STOMP (1/6)

Scalable Time series Ordered Matrix Profile

$O(n^2)$ time, $O(n)$ space algorithm to compute the matrix profile.

Matrix profile algorithm – STOMP (2/6)

Scalable Time series Ordered Matrix Profile

Z-normalized Euclidean distance
 m := number of data points

$$T_i T_j = \sum_{k=0}^{m-1} t_{i+k} t_{j+k} \quad \text{Dot product of the } i^{\text{th}} \text{ window and the } j^{\text{th}} \text{ window.}$$

$$d_{i,j} = \sqrt{2m \left(1 - \frac{T_i T_j - m\mu_i \mu_j}{m\sigma_i \sigma_j} \right)}$$

- We can precompute and store the means and standard deviations in $O(n)$ space and time
- Once we know T_i, T_j , it takes $O(1)$ time to compute the distance $d_{i,j}$

Matrix profile algorithm – STOMP (3/6)

Scalable Time series Ordered Matrix Profile

- The relationship between $T_i T_j$ and $T_{i+1} T_{j+1}$

$$T_i T_j = \sum_{k=0}^{m-1} t_{i+k} t_{j+k}$$

| | | | | | | | | | |
|-----|----------|-----------|-----------|----------|-------------|-----------|-----|----------|--|
| ... | t_i | t_{i+1} | t_{i+2} | ... | t_{i+m-1} | t_{i+m} | ... | | |
| | \times | \times | $+$ | \times | $+$ | \dots | $+$ | \times | |
| ... | t_j | t_{j+1} | t_{j+2} | ... | t_{j+m-1} | t_{j+m} | ... | | |

$$T_{i+1} T_{j+1} =$$

| | | | | | | | | | |
|-----|-------|-----------|-----------|----------|-------------|-----------|-----|----------|--|
| ... | t_i | t_{i+1} | t_{i+2} | ... | t_{i+m-1} | t_{i+m} | ... | | |
| | | \times | $+$ | \times | $+$ | \dots | $+$ | \times | |
| ... | t_j | t_{j+1} | t_{j+2} | ... | t_{j+m-1} | t_{j+m} | ... | | |

$$T_{i+1} T_{j+1} = T_i T_j - t_i t_j + t_{i+m} t_{j+m}$$

$O(1)$ time complexity

Matrix profile algorithm – STOMP (4/6)

Scalable Time series Ordered Matrix Profile

- 1) All means and standard deviations are precomputed. This costs linear time and space.

| | | | | |
|------------|------------|------------|-----|------------------|
| μ_1 | μ_2 | μ_3 | ... | μ_{n-m+1} |
| σ_1 | σ_2 | σ_3 | ... | σ_{n-m+1} |

Matrix profile algorithm – STOMP (5/6)

Scalable Time series Ordered Matrix Profile

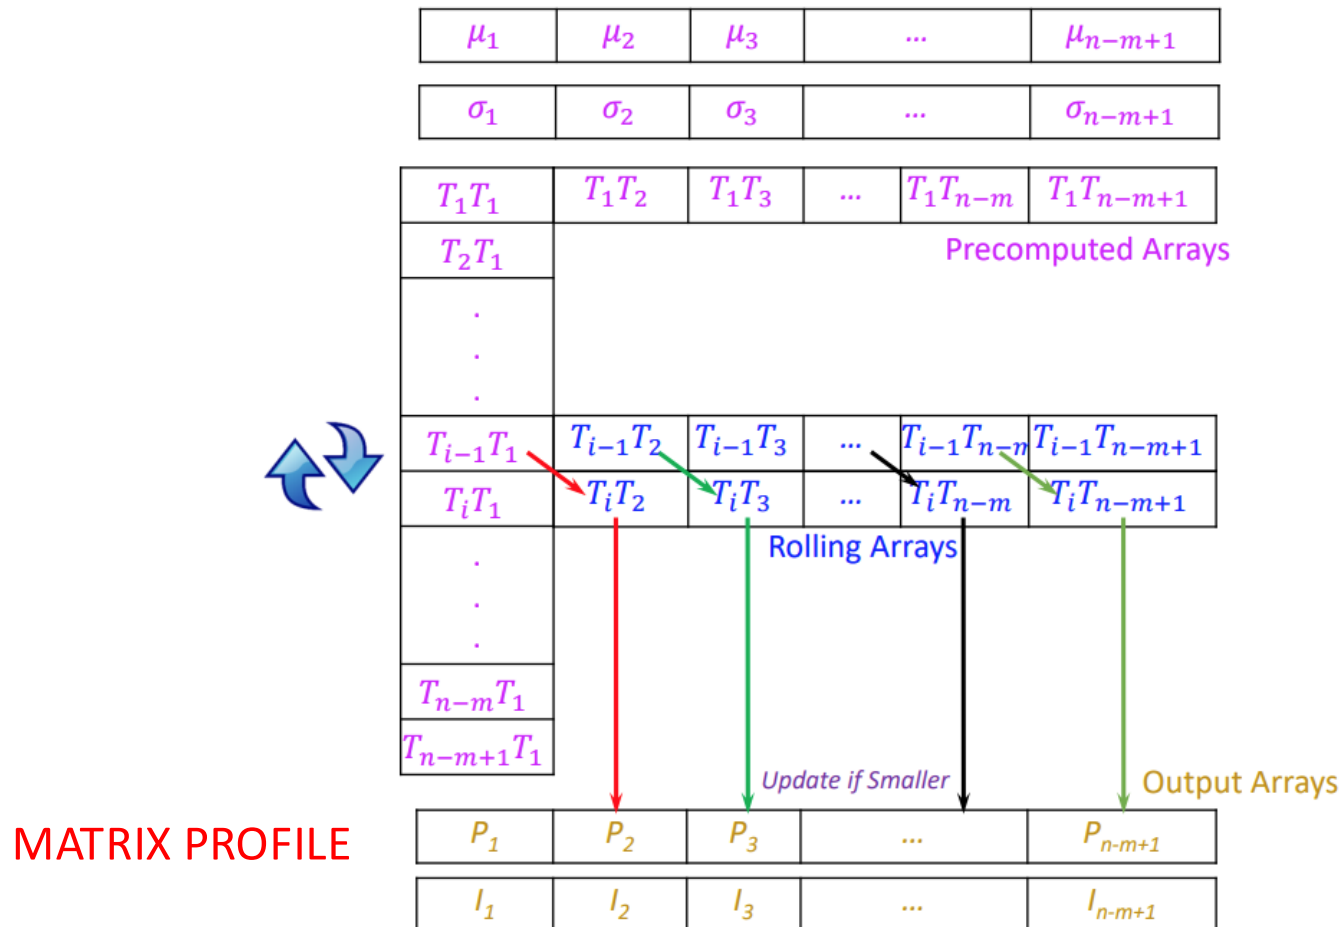
2) The first column and row of the matrix are identical and pre-computed by MASS.

| | | | | | |
|------------|--------------------|------------|-----|------------------|----------------|
| μ_1 | μ_2 | μ_3 | ... | μ_{n-m+1} | |
| σ_1 | σ_2 | σ_3 | ... | σ_{n-m+1} | |
| T_1T_1 | T_1T_2 | T_1T_3 | ... | T_1T_{n-m} | T_1T_{n-m+1} |
| T_2T_1 | Precomputed Arrays | | | | |

Matrix profile algorithm – STOMP (6/6)

Scalable Time series Ordered Matrix Profile

3) The algorithm iterates through the rows. The previous row is maintained as a local array to feed dot products.



Notebook motif discovery

- <https://github.com/target/matrixprofile-ts/blob/master/docs/examples/Motif%20Discovery.ipynb>

- Matrix Profile library:

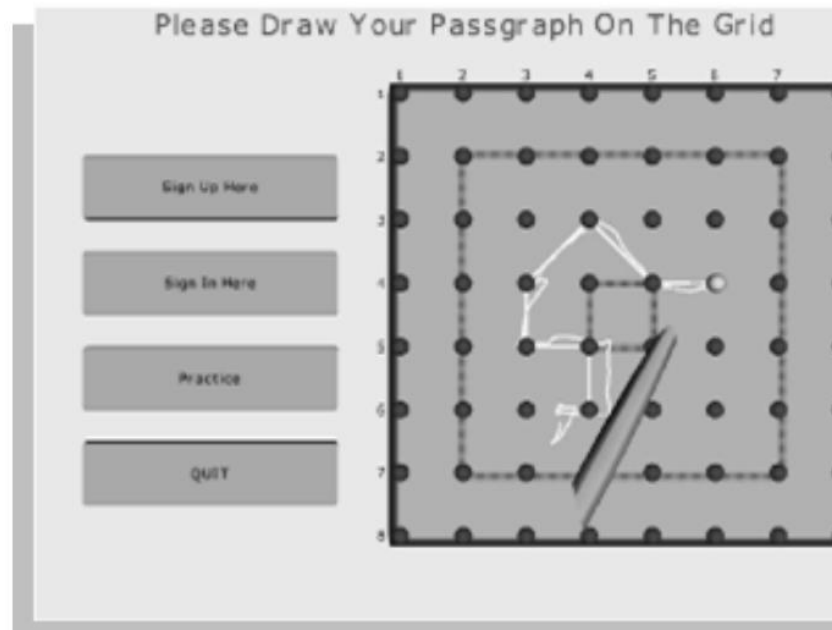
<https://github.com/target/matrixprofile-ts/tree/master>



Presentation du TD

Haptics DATA

Data are taken from 5 people entering their passgraph (a code to access a system protected by a graphical authentication system) on a touchscreen. The data are the x-axis movement only.



Assignment

- Open the Python Notebook

TS_MotifDiscovery.ipynb

Load the modules in the folder

- Instruction are contained in the notebook

References

- Time Series Data Mining Using the Matrix Profile: A Unifying View of Motif Discovery, Anomaly Detection, Segmentation, Classification, Clustering and Similarity Joins www.cs.ucr.edu/~eamonn/MatrixProfile.html
- [https://stumpy.readthedocs.io/en/latest/Tutorial The Matrix Profile.html](https://stumpy.readthedocs.io/en/latest/Tutorial%20The%20Matrix%20Profile.html)
- Matrix Profile I: All Pairs Similarity Joins for Time Series: A Unifying View That Includes Motifs, Discords and Shapelets
<https://ieeexplore.ieee.org/document/7837992>
- <https://github.com/matrix-profile-foundation>